



UCL

Information Retrieval & Data Mining [COMP0084]

Introduction to machine learning & data mining – Part 1

Vasileios Lamos

Computer Science, UCL

@lamos



- ▶ In this lecture:
 - Association rule mining (*data mining*) – Apriori algorithm
 - Introduction to machine learning – Part 1
- ▶ Useful additional reads
 - Chapters 2 and 4 of “*Web Data Mining*” by Bing Liu (2006)
 - Chapters 3 and 4 of “*The Elements of Statistical Learning*” by Hastie, Tibshirani, and Friedman (2008)
 - Chapter 5 of “*Speech and language processing*” by Jurafsky and Martin (2021)
- ▶ Some slides adapted from Bing Liu’s course — cs.uic.edu/~liub/teach/cs583-fall-21/cs583.html
- ▶ Many slides were adapted from Prof. Emine Yilmaz’s lectures in previous years

Data mining — Definition

- ▶ **Data mining** is the process of discovering (*mining*) useful patterns from or conducting inferences based on various types of *data* sources such as structured information repositories (e.g. databases), text, images, sound, video, and so on.
- ▶ **Multi-disciplinary**: machine learning (or AI more broadly), statistics, databases, information retrieval — *but the distinction between machine learning and data mining is becoming increasingly difficult, especially from an applications perspective.*
- ▶ Strong research community: Knowledge Discovery and Data Mining or **KDD** — kdd.org
- ▶ Why? Gaining knowledge from a database is not as simple issue database queries
- ▶ Applications include marketing, recommendations, scientific data analysis, and *any task involving large amounts of data*

- ▶ Today: a basic look into **Association rule mining / learning** – *perhaps the most important task proposed and studied by the data mining community*
- ▶ Introduced by Agrawal, Imielinski, and Swami in 1993 – dl.acm.org/doi/pdf/10.1145/170035.170072
- ▶ Applicable on categorical / discrete data (e.g. product categories, movies, songs)
- ▶ Initially used for market basket analysis to understand how products purchased by customers are related, e.g.

spaghetti → *basil* [support = 0.1%, confidence = 25%]

Association rule mining – Notation & definitions

market basket
transactions

$t_1 : \{\text{almonds, cashews, pistachios}\}$
 $t_2 : \{\text{almonds, bananas}\}$
...
 $t_n : \{\text{cashews, oranges, pistachios}\}$

- ▶ A set of all the m **items**, $I = \{i_1, i_2, \dots, i_m\}$
– e.g. “almonds” is an item
- ▶ A set of all the n **transactions**, $T = \{t_1, t_2, \dots, t_n\}$
- ▶ A transaction t_i is a set of items, and hence $t_i \subseteq I$

Association rule mining – Notation & definitions

market basket
transactions

$$\begin{aligned} t_1 &: \{\text{almonds, cashews, pistachios}\} \\ t_2 &: \{\text{almonds, bananas}\} \\ &\dots \\ t_n &: \{\text{cashews, oranges, pistachios}\} \end{aligned}$$

- ▶ An **itemset** is a set of items
 - e.g. $X = \{\text{almonds, cashews}\}$
- ▶ A **k -itemset** is an itemset with k items
 - e.g. $X = \{\text{almonds, cashews, pistachios}\}$ is a 3-itemset
- ▶ A transaction t_i contains the set of items (**itemset**) $X \subseteq I$, if $X \subseteq t_i$
- ▶ An **association rule** between itemsets X, Y is an implication of the form:

$$X \rightarrow Y, \text{ where } X, Y \subset I, \text{ and } X \cap Y = \emptyset$$

- ▶ **Association rule** (a *pattern*): $X \rightarrow Y$

- when X occurs, Y occurs with a certain *support* and *confidence*

- ▶ **support** =
$$\frac{(X \cup Y) . \text{count}}{n}$$

- probability that a transaction will contain both itemsets X and Y , $\Pr(X \cup Y)$

- how many times X and Y appear together in all (n) transactions in T divided by n

- ▶ Association rule (a *pattern*): $X \rightarrow Y$
 - when X occurs, Y occurs with a certain *support* and *confidence*
- ▶ **support** = $\frac{(X \cup Y) . \text{count}}{n} \sim \text{Pr}(X \cup Y)$
- ▶ **confidence** = $\frac{(X \cup Y) . \text{count}}{X . \text{count}}$
 - conditional probability that a transaction that contains X will also contain Y , $\text{Pr}(Y | X)$
 - how many times a transaction that contains X also contains Y divided by the number of transactions that contain X

- ▶ **Goal:** Find all association rules ($X \rightarrow Y$) that satisfy a pre-specified (*by us!*) **minimum support** (also abbreviated as **minsup**) and **minimum confidence** (**minconf**)
- ▶ Key features
 - **Completeness**, i.e. find all rules
Note that $X \rightarrow Y$ and $Y \rightarrow X$ are different rules. **Why?**
 - Mining with data on hard disk (*because it is not always feasible to load everything in memory*)

Association rule mining – An example

- ▶ Transactions

- ▶ Let's set

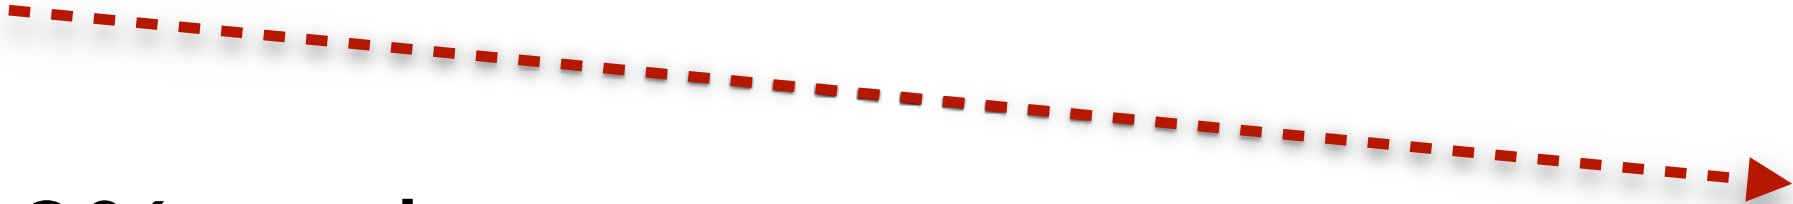
- **minsup** = 30%, and
- **minconf** = 80%

- ▶ **Frequent itemset** examples:

- {almonds, cashews} with support 3/7 (> **minsup**)
- {cashews, pistachios} with support 4/7
- {cashews, oranges, pistachios} with support 3/7

- ▶ **Association rule candidates** from the above frequent itemsets

- almonds → cashews with confidence 3/4 (< **minconf**, *rejected*)
- pistachios → cashews with confidence 4/4 (> **minconf**, *accepted*)
- {cashews, oranges} → pistachios with confidence 3/3 (*accepted*)



t_1 : {almonds, cashews, pistachios}
 t_2 : {almonds, bananas}
 t_3 : {apples, bananas}
 t_4 : {almonds, bananas, cashews}
 t_5 : {almonds, bananas, cashews, oranges, pistachios}
 t_6 : {cashews, oranges, pistachios}
 t_7 : {cashews, oranges, pistachios}

- ▶ Large number of different association rule mining algorithms
- ▶ Different strategies, data structures, computational efficiency, memory requirements
- ▶ But their output can only be the same:
 - Given a transaction data set T , minsup, and minconf, the set of association rules in T is uniquely determined.
- ▶ Let's briefly look at a foundational algorithm for association rule mining: **Apriori**

- ▶ **Apriori** is perhaps the most popular algorithm in data mining
- ▶ “Apriori” probably because it uses “prior” knowledge of frequent itemsets
- ▶ Proposed by Agrawal and Srikant in 1994 — vldb.org/conf/1994/P487.pdf
- ▶ Same two steps (*that we’ve seen previously*)
 - find all the itemsets with a minimum support (frequent itemsets)
 - then use the frequent itemsets to generate association rules

Apriori — Identify frequent itemsets

- ▶ The key idea of Apriori is the downward closure property (also known as the “Apriori property”):
 - Any subset of a frequent itemset is also a frequent itemset
 - = Any subset of an itemset whose support is $\geq \mathbf{minsup}$ has also support that is $\geq \mathbf{minsup}$
- ▶ If the itemset $\{a, b, c, d\}$ with 4 items is frequent, then the $(2^4 - 2)$ non-empty sub-itemsets will also be frequent. These are: $\{a\}$, $\{b\}$, $\{c\}$, $\{d\}$, $\{a, b\}$, $\{a, c\}$, $\{a, d\}$, $\{b, c\}$, $\{b, d\}$, $\{c, d\}$, $\{a, b, c\}$, $\{a, b, d\}$, $\{a, c, d\}$, and $\{b, c, d\}$.
- ▶ **Contraposition:** if an itemset is not frequent, then any of its supersets cannot be frequent

Apriori — The gist of the algorithm

- ▶ Apriori is an iterative algorithm
 - given a minimum support
 - find all frequent 1-itemsets (denoted by $F[1]$ in the source code)
 - use those to find all frequent 2-itemsets, and so on
 - > $C[2]$ is a list of frequent 2-itemset candidates based on $F[1]$
 - > $F[2] \subseteq C[2]$ is a list with the frequent 2-itemsets
 - in each iteration k of the algorithm only consider itemsets that contain some frequent $(k - 1)$ -itemset

Apriori — An important detail (*item ordering*)

- ▶ Items should be sorted according to a sorting scheme — i.e. lexicographic order
- ▶ This order will be used throughout the algorithm as it helps to reduce redundant passes on the data, e.g. the frequent itemset {a, b, c, d} is identical to the frequent itemsets {c, d, a, b} or {b, a, d, c} — we only need to deal with {a, b, c, d} once.

Apriori — Pseudocode of the algorithm (*part 1*)

```
01 % T: all the transactions, MINSUP: frequent itemset minimum support
02 function apriori(T, MINSUP) :
03     % C[1] count of 1-itemsets, n transactions in T
04     C[1], n ← initial-pass(T)
05     % F[1] is the set of frequent 1-itemsets
06     F[1] ← {f | f in C[1] AND f.count/n ≥ MINSUP}
07     for k = 2; F[k-1] ≠ ∅; k++:
08         % use the (k-1)-itemsets to generate k-itemset candidates, C[k]
09         C[k] ← generate-candidates(F[k-1])
10         for each transaction t in T:
11             for each candidate c in C[k]:
12                 if c is in t:
13                     c.count++
14         F[k] ← {c in C[k] | c.count/n ≥ MINSUP}
15
16     return F
```

- ▶ The **generate-candidates** function takes the $(k - 1)$ -frequent itemsets, denoted by $F[k-1]$ in the source code, and returns a superset of k -frequent itemset candidates, denoted by $C[k]$
- ▶ Two steps
 - **Join**: generate all possible candidate k -itemsets $C[k]$ based on $F[k-1]$
 - **Prune**: remove those candidates in $C[k]$ that cannot be frequent, i.e. if a candidate itemset has a subset of items that is not already identified as a frequent itemset it should be removed

Apriori — Pseudocode of the algorithm (part 2)

```
01 % using frequent (k-1)-itemsets generate frequent k-itemset candidates
02 function generate-candidates (F[k-1]) :
03     C[k] ← ∅
04     for every f1, f2 in F[k-1] where:
05         a = f1 - f2 AND                                % set difference
06         b = f2 - f1 AND                                % set difference
07         (a AND b) are both of size 1 AND                % f1 and f2 differ by 1 element
08         a < b do:                                       % lexicographic comparison
09             c ← {f1, b}                                  % frequent k-itemset candidate
10             C[k] ← {C[k], c}
11             for each (k-1)-subset s of c do:
12                 if s not in F[k-1]:
13                     delete c from C[k]                  % pruning non-frequent candidates
14
15     return C[k]
```

Apriori — An example

`t[1]: {almonds, cashews, pistachios}`

`t[2]: {almonds, bananas}`

`t[3]: {apples, bananas}`

`t[4]: {almonds, bananas, cashews}`

`t[5]: {almonds, bananas, cashews, oranges, pistachios}`

`t[6]: {cashews, oranges, pistachios}`

`t[7]: {cashews, oranges, pistachios}`

Let's use Apriori to identify all frequent itemsets with minimum support of 30%

Apriori — An example

```
t[1]: {almonds, cashews, pistachios}
t[2]: {almonds, bananas}
t[3]: {apples, bananas}
t[4]: {almonds, bananas, cashews}
t[5]: {almonds, bananas, cashews, oranges, pistachios}
t[6]: {cashews, oranges, pistachios}
t[7]: {cashews, oranges, pistachios}
```

C[1]: {almonds:4/7, apples:1/7, bananas:4/7, cashews:5/7, oranges:3/7, pistachios:4/7}

F[1]: {almonds, bananas, cashews, oranges, pistachios}

C[2]: { {almonds, bananas}:3/7, {almonds, cashews}:3/7,
 {almonds, oranges}:1/7, {almonds, pistachios}:2/7,
 {bananas, cashews}:2/7, {bananas, oranges}:1/7,
 {bananas, pistachios}:1/7, {cashews, oranges}:3/7,
 {cashews, pistachios}:4/7, {oranges, pistachios}:3/7 }

Apriori — An example

```
t[1]: {almonds, cashews, pistachios}
t[2]: {almonds, bananas}
t[3]: {apples, bananas}
t[4]: {almonds, bananas, cashews}
t[5]: {almonds, bananas, cashews, oranges, pistachios}
t[6]: {cashews, oranges, pistachios}
t[7]: {cashews, oranges, pistachios}
```

```
C[2]: { {almonds, bananas}:3/7,      {almonds, cashews}:3/7,
        {almonds, oranges}:1/7,     {almonds, pistachios}:2/7,
        {bananas, cashews}:2/7,     {bananas, oranges}:1/7,
        {bananas, pistachios}:1/7,  {cashews, oranges}:3/7,
        {cashews, pistachios}:4/7,  {oranges, pistachios}:3/7 }
```

```
F[2]: { {almonds, bananas}, {almonds, cashews}, {cashews, oranges},
        {cashews, pistachios}, {oranges, pistachios} }
```

Apriori – An example

```
t[1]: {almonds, cashews, pistachios}
t[2]: {almonds, bananas}
t[3]: {apples, bananas}
t[4]: {almonds, bananas, cashews}
t[5]: {almonds, bananas, cashews, oranges, pistachios}
t[6]: {cashews, oranges, pistachios}
t[7]: {cashews, oranges, pistachios}
```

```
F[2]: { {almonds, bananas}, {almonds, cashews}, {cashews, oranges},
        {cashews, pistachios}, {oranges, pistachios} }
```

```
C[3]: { {almonds, bananas, cashews}:2/7,
        {cashews, oranges, pistachios}:3/7 } }
```

*** Incorrect ***

```
C[3]: { {cashews, oranges, pistachios}:3/7 }
```

entry {almonds, bananas, cashews} **will be pruned because**
{bananas, cashews} **is not in** F[2]

```
F[3]: { {cashews, oranges, pistachios} }
```

Apriori — An example

```
t[1]: {almonds, cashews, pistachios}
t[2]: {almonds, bananas}
t[3]: {apples, bananas}
t[4]: {almonds, bananas, cashews}
t[5]: {almonds, bananas, cashews, oranges, pistachios}
t[6]: {cashews, oranges, pistachios}
t[7]: {cashews, oranges, pistachios}
```

Apriori identified the following *frequent* itemsets with a minimum support of 30%:

$F[1]: \{\text{almonds}:4/7, \text{bananas}:4/7, \text{cashews}:5/7, \text{oranges}:3/7, \text{pistachios}:4/7\}$

$F[2]: \{ \{\text{almonds}, \text{bananas}\}:3/7, \quad \{\text{almonds}, \text{cashews}\}:3/7,$
 $\quad \{\text{cashews}, \text{oranges}\}:3/7, \quad \{\text{cashews}, \text{pistachios}\}:4/7,$
 $\quad \{\text{oranges}, \text{pistachios}\}:3/7 \}$

$F[3]: \{ \{\text{cashews}, \text{oranges}, \text{pistachios}\}:3/7 \}$

Apriori — Generating association rules from frequent itemsets

- ▶ Frequent itemsets do not directly provide association rules
- ▶ For each frequent itemset F
For each non-empty subset A of F (*no repetitions*)
 - $B = F - A$
 - $A \rightarrow B$ is an association rule if confidence $(A \rightarrow B) \geq \mathbf{minconf}$
 $\text{support}(A \rightarrow B) = \text{support}(A \cup B) = \text{support}(F)$
$$\text{confidence}(A \rightarrow B) = \frac{\text{support}(A \cup B)}{\text{support}(A)}$$

Apriori – Generating association rules (*example*)

```
t[1]: {almonds, cashews, pistachios}
t[2]: {almonds, bananas}
t[3]: {apples, bananas}
t[4]: {almonds, bananas, cashews}
t[5]: {almonds, bananas, cashews, oranges, pistachios}
t[6]: {cashews, oranges, pistachios}
t[7]: {cashews, oranges, pistachios}
```

minsup = 30%, minconf = 80%, let's use $F[3] : \{ \{ \text{cashews, oranges, pistachios} \} : 3/7 \}$

$A = \{ \{ \text{cashews, oranges} \}, \{ \text{cashews, pistachios} \}, \{ \text{oranges, pistachios} \}, \{ \text{cashews} \}, \{ \text{oranges} \}, \{ \text{pistachios} \} \}$

$A \rightarrow B$

$\{ \text{cashews, oranges} \}$	\rightarrow pistachios	confidence = 1
$\{ \text{cashews, pistachios} \}$	\rightarrow oranges	confidence = 0.75
$\{ \text{oranges, pistachios} \}$	\rightarrow cashews	confidence = 1
cashews	\rightarrow {oranges, pistachios}	confidence = 0.6
oranges	\rightarrow {cashews, pistachios}	confidence = 1
pistachios	\rightarrow {cashews, oranges}	confidence = 0.75

Apriori – Generating association rules (*example*)

```
t[1]: {almonds, cashews, pistachios}
t[2]: {almonds, bananas}
t[3]: {apples, bananas}
t[4]: {almonds, bananas, cashews}
t[5]: {almonds, bananas, cashews, oranges, pistachios}
t[6]: {cashews, oranges, pistachios}
t[7]: {cashews, oranges, pistachios}
```

minsup = 30%, **minconf** = 80%, let's use $F[3] : \{ \{ \text{cashews, oranges, pistachios} \} : 3/7 \}$

$A = \{ \{ \text{cashews, oranges} \}, \{ \text{cashews, pistachios} \}, \{ \text{oranges, pistachios} \}, \{ \text{cashews} \}, \{ \text{oranges} \}, \{ \text{pistachios} \} \}$

A → B

{cashews, oranges}	→ pistachios	confidence = 1
{cashews, pistachios}	→ oranges	confidence = 0.75
{oranges, pistachios}	→ cashews	confidence = 1
cashews	→ {oranges, pistachios}	confidence = 0.6
oranges	→ {cashews, pistachios}	confidence = 1
pistachios	→ {cashews, oranges}	confidence = 0.75

- ▶ To obtain an association rule $A \rightarrow B$, we need to compute the quantities: support ($A \cup B$) and support (A)
- ▶ This information has already been recorded during itemset generation. No need to access the raw transaction data any longer.
- ▶ Not as time consuming as frequent itemset generation, although there are efficient algorithms to generate association rules as well

The (*very*) basics of machine learning

- definition
- supervised learning (*regression, classification*)
- unsupervised learning

Machine learning

- ▶ Arthur Samuel (IBM, 1959): “*Machine learning is the field of study that gives the computer the ability to learn (a task) without being explicitly programmed.*”
 - credited for coining the term
 - although we are still explicitly programming them to learn!
- ▶ Tom Mitchell (CMU, 1998): “A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E .”
 - more formal definition
 - learning from experience (observations, data)



Notational conventions for this lecture

$x \in \mathbb{R}$ denotes a real-valued scalar

$\mathbf{x} \in \mathbb{R}^n$ denotes a real-value vector with n elements

$\mathbf{X} \in \mathbb{R}^{n \times m}$ denotes a real-valued matrix with n rows and m columns

$\mathbf{y} \in \mathbb{R}^m$ denotes m instances of a real valued response (or target) variable

$\hat{\mathbf{y}} \in \mathbb{R}^m$ denotes m inferences of a real valued response variable

$\|\mathbf{x}\|_k = \left(\sum_{i=1}^n |x_i|^k \right)^{\frac{1}{k}}$ denotes the L_k -norm of $\mathbf{x} \in \mathbb{R}^n$

Learning from experience

- ▶ Experience is something tangible, i.e. an observation and eventually a data point, something that can take a numeric form
- ▶ \mathbf{x}_i denotes a numeric interpretation of an input
 y_i denotes a numeric interpretation of an output

$\langle \mathbf{x}_i, y_i \rangle$ is an observation / sample

