



UCL

Information Retrieval & Data Mining [COMP0084]

Introduction to machine learning & data mining – Part 2

Vasileios Lamos

Computer Science, UCL

@lamos



Preliminaries

- ▶ In this lecture:
 - Introduction to machine learning – Part 2
 - Supervised learning: regression, classification
 - Unsupervised learning: clustering
 - With examples!
- ▶ Useful additional reads
 - Chapters 2, 4 of “*Web Data Mining*” by Bing Liu (2006)
 - Chapters 3, 4, 14 of “*The Elements of Statistical Learning*” by Hastie, Tibshirani, and Friedman (2008)
 - Chapter 5 of “*Speech and language processing*” by Jurafsky and Martin (2021)
 - *Advanced* – Flu rate prediction from web search activity, [nature.com/articles/srep12760.pdf](https://www.nature.com/articles/srep12760.pdf)
- ▶ Some slides adapted from Bing Liu’s course — cs.uic.edu/~liub/teach/cs583-fall-21/cs583.html
- ▶ Many slides were adapted from Prof. Emine Yilmaz’s lectures in previous years

Machine learning

- ▶ Arthur Samuel (IBM, 1959): “*Machine learning is the field of study that gives the computer the ability to learn (a task) without being explicitly programmed.*”
 - credited for coining the term
 - although we are still explicitly programming them to learn!
- ▶ Tom Mitchell (CMU, 1998): “A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E .”
 - more formal definition
 - learning from experience (observations, data)



Notational conventions for this lecture

$x \in \mathbb{R}$ denotes a real-valued scalar

$\mathbf{x} \in \mathbb{R}^n$ denotes a real-value vector with n elements

$\mathbf{X} \in \mathbb{R}^{n \times m}$ denotes a real-valued matrix with n rows and m columns

$\mathbf{y} \in \mathbb{R}^m$ denotes m instances of a real valued response (or target) variable


$\hat{\mathbf{y}} \in \mathbb{R}^m$ denotes m inferences of a real valued response variable

$\|\mathbf{x}\|_p = \left(\sum_{i=1}^n |x_i|^p \right)^{\frac{1}{p}}$ denotes the L_p -norm of $\mathbf{x} \in \mathbb{R}^n, p \in \mathbb{N}_{>0}$

Learning from experience

- ▶ Experience is something tangible, i.e. an observation and eventually a data point, something that can take a numeric form
- ▶ \mathbf{x}_i denotes a numeric interpretation of an input
 y_i denotes a numeric interpretation of an output

$\langle \mathbf{x}_i, y_i \rangle$ is an observation / sample



Common machine learning categorisation

► Supervised learning

Learn a mapping f from inputs \mathbf{X} to outputs \mathbf{y} – also can be expressed by $f : \mathbf{X} \rightarrow \mathbf{y}$

– \mathbf{X} are also called features, observations, covariates, predictors

– \mathbf{y} are also called labels, targets, responses, ground truth

– $\langle \mathbf{X}, \mathbf{y} \rangle$ can also be referred to as observations or samples

► Unsupervised learning

No outputs associated with the input \mathbf{X} – the task becomes to discover an underlying structure or patterns in \mathbf{X}

► Reinforcement learning

The system or agent has to learn how to interact with its environment

Policy: which action to take in response to an input \mathbf{X}

Different from supervised learning because no definitive responses are given

Only rewards – *learning with a critic as opposed to learning with a teacher*

► Regression

estimate / predict a continuous output / target variable

i.e. learn $f: \mathbf{X} \in \mathbb{R}^{n \times m} \rightarrow \mathbf{y} \in \mathbb{R}^n$

Examples: predict a time series trend (e.g. in finance), estimate the prevalence of a condition in epidemiology

► Classification


estimate a set of C unordered (and mutually exclusive) labels / classes

i.e. learn $f: \mathbf{X} \in \mathbb{R}^{n \times m} \rightarrow \mathbf{y} \in \{1, 2, \dots, C\}$

Examples: detect spam email, medical imaging, text/document classification


Supervised learning – Regression

- Estimate the **prevalence of influenza-like illness** in England based on the **frequency of the search query** “*how long does flu last*”




Supervised learning – Regression

- ▶ Estimate the **prevalence of influenza-like illness** in England based on the **frequency of the search query** “*how long does flu last*”
- ▶ Linearly related, bivariate correlation of 0.975
- ▶ Can we capture this relationship with a straight line?
- ▶ **Data:** dropbox.com/s/rgyg190whw26qrj/data-COMP0084-intro-to-ml.zip?dl=0




Supervised learning – Regression

- ▶ Estimate the prevalence of influenza-like illness in England based on the frequency of the search query “*how long does flu last*”
- ▶ Linearly related, bivariate correlation of 0.975
- ▶ Can we capture this relationship with a straight line?



Supervised learning – Regression

- ▶ Estimate the prevalence of influenza-like illness in England based on the frequency of the search query “*how long does flu last*”
- ▶ Linearly related, bivariate correlation of 0.975
- ▶ Can we capture this relationship with a straight line?
- ▶ Which line is the “best” though?




Supervised learning — Ordinary least squares (*linear*) regression

- ▶ \mathbf{y} denotes the weekly influenza-like illness prevalence in England from September 2017 until the end of August 2018
- ▶ \mathbf{x} denotes the corresponding weekly frequency of the search query “how long does flu last” (Google) for the same time period
- ▶ We want to learn a linear mapping f from the input \mathbf{x} to the output \mathbf{y} based on our current observations, i.e. for a weekly query frequency x_i , $f(x_i) = \hat{y}_i = \alpha x_i + \beta \approx y_i$
- ▶ This linear mapping has two unknown hyper-parameters: $\{\alpha, \beta\}$
- ▶ Find a line that best fits to our observations


Supervised learning – Ordinary least squares (*linear*) regression

- ▶ $\mathbf{y} \sim$ weekly flu prevalence
- ▶ $\mathbf{x} \sim$ weekly search frequency of “*how long does flu last*”
- ▶ $f: \mathbf{x} \rightarrow \mathbf{y}$ such that
 $f(x_i) = \hat{y}_i = \alpha x_i + \beta \approx y_i$
- ▶ Find a line that best fits to our observations using **ordinary least squares (OLS)** regression




Supervised learning – OLS regression, *alternative point of view*

- ▶ $\mathbf{y} \sim$ weekly flu prevalence
- ▶ $\mathbf{x} \sim$ weekly search frequency of “*how long does flu last*”
- ▶ $f: \mathbf{x} \rightarrow \mathbf{y}$ such that
 $f(x_i) = \hat{y}_i = \alpha x_i + \beta \approx y_i$
- ▶ Find a line that best fits to our observations using **ordinary least squares (OLS)** regression




Supervised learning – OLS regression, *alternative point of view*

- ▶ $\mathbf{y} \sim$ weekly flu prevalence
- ▶ $\mathbf{x} \sim$ weekly search frequency of “*how long does flu last*”
- ▶ $f: \mathbf{x} \rightarrow \mathbf{y}$ such that
 $f(x_i) = \hat{y}_i = \alpha x_i + \beta \approx y_i$
- ▶ Find a line that best fits to our observations using **ordinary least squares (OLS)** regression




Supervised learning – OLS regression, *alternative point of view*

- ▶ $\mathbf{y} \sim$ weekly flu prevalence
- ▶ $\mathbf{x} \sim$ weekly search frequency of “*how long does flu last*”
- ▶ $f: \mathbf{x} \rightarrow \mathbf{y}$ such that $f(x_i) = \hat{y}_i = \alpha x_i + \beta \approx y_i$
- ▶ Find a line that best fits to our observations using **ordinary least squares (OLS)** regression



Supervised learning – OLS regression, *alternative point of view*

- ▶ $\mathbf{y} \sim$ weekly flu prevalence
- ▶ $\mathbf{x} \sim$ weekly search frequency of “*how long does flu last*”
- ▶ $f: \mathbf{x} \rightarrow \mathbf{y}$ such that $f(x_i) = \hat{y}_i = \alpha x_i + \beta \approx y_i$
- ▶ Find a line that best fits to our observations using **ordinary least squares (OLS)** regression



Supervised learning – OLS regression calculus solution

- ▶ The aim is to learn $f: \mathbf{X} \in \mathbb{R}^{n \times m} \rightarrow \mathbf{y} \in \mathbb{R}^n$
- ▶ f is a linear function, a set of weights and an intercept term; denoted by $\mathbf{w} \in \mathbb{R}^m$
- ▶ In our regression task (*see previous slides*), there is 1 weight (α) and the intercept (β)
 \mathbf{X} has one column with the values of \mathbf{x} and the other column is 1s
- ▶ Minimise a loss function known as residual sum or squares (*equivalent to mean squared error that we will see next*): $\mathcal{L}(\mathbf{w}) = \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 = (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$
- ▶ This can also be written as: $\mathcal{L}(\mathbf{w}) = \mathcal{L}(\alpha, \beta) = \sum_{i=1}^n (\alpha x_i + \beta - y_i)^2$


Supervised learning – OLS regression calculus solution

- ▶ The aim is to learn $f: \mathbf{X} \in \mathbb{R}^{n \times m} \rightarrow \mathbf{y} \in \mathbb{R}^n$
- ▶ f is a linear function, a set of weights and an intercept term; denoted by $\mathbf{w} \in \mathbb{R}^m$
- ▶ In our regression task (*see previous slides*), there is 1 weight (α) and the intercept (β)
 \mathbf{X} has one column with the values of \mathbf{x} and the other column is 1s
- ▶ Minimise a loss function known as residual sum or squares (*equivalent to mean squared error that we will see next*): $\mathcal{L}(\mathbf{w}) = \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 = (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$
- ▶ Derivative with respect to \mathbf{w} : $\frac{\partial \mathcal{L}}{\partial \mathbf{w}} = -2\mathbf{X}^\top \mathbf{y} + 2\mathbf{X}^\top \mathbf{X}\mathbf{w}$
- ▶ Set this to 0 and hence $\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$ as long as $\mathbf{X}^\top \mathbf{X}$ is full rank which means that the observations (rows) in \mathbf{X} are more than the features ($n > m$) and that the features have no linear dependence

- ▶ Going back to our example, $\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$ would give $\mathbf{w} = [0.93351 \ -0.036631]$, i.e. $\alpha = 0.93351$ and $\beta = -0.036631$
- ▶ The question now becomes, **how well will this model do in the next flu season?**
- ▶ Let's use the above values of α and β to estimate weekly flu prevalence in England for the season 2018/19 based on the corresponding frequency of the search query "*how long does flu last*"
- ▶ And then compare it with the actual flu prevalence in England for 2018/19


Supervised learning – OLS model training & testing

- ▶ These (**red** line, dot • marker) are the estimated (*inferred*) flu rates in 2018/19 (to be exact from *September 2018 to August 2019*) based on the OLS model and the frequency of the search query “*how long does flu last*”
- ▶ Recall, we trained our model using non-overlapping data from 2017/18 (*September 2017 to August 2018*)



Supervised learning – OLS model training & testing

- ▶ The **black solid** line represents the corresponding flu rates as reported by a health agency in the UK
- ▶ *Do you think this simple OLS model based on a single web search query did well?*
- ▶ $r = 0.919$ (bivariate correlation)
RMSE = 0.0632 (root mean squared error)
MAE = 0.0519 (mean absolute error)
- ▶ considering the simplicity of the model, its accuracy is quite surprising



- ▶ **Gradient descent:** optimisation algorithm that minimises a loss function \mathcal{J} with respect to a set of hyperparameters
- ▶ Loss function for ordinary least squares (OLS) regression? If $\hat{\mathbf{y}} = \mathbf{X}\mathbf{w}$ denotes our estimates for \mathbf{y} , then the loss function for OLS is their mean squared difference (error):

$$\mathcal{J}(\mathbf{w}) = \frac{1}{2n} \sum_{i=1}^n (\hat{y}_i - y_i)^2, \text{ where } \hat{y}_i \in \hat{\mathbf{y}}, y_i \in \mathbf{y}$$

- ▶ **Basic steps** of gradient descent
 - define a loss function, \mathcal{J}
 - compute the partial derivatives of \mathcal{J} w.r.t. each hyperparameter
 - update hyperparameters using their partial derivatives and learning rate ℓ often $\in (0,1)$
 - repeat until convergence

- ▶ **Learning rate:** how far away are we going to go in the opposite direction of the partial derivative — *we are going to see an example of this*
- ▶ **Why does it work?** We are taking steps in the opposite direction of the partial gradient to identify a local minimum.
- ▶ **When does it not work?** Not directly applicable to non-differentiable loss functions (but there exist workarounds)

In our example, we are modelling a flu rate y_i using the frequency of a search query x_i

- ▶ **Hypothesis:** $\hat{y}_i = \alpha x_i + \beta$
— a flu estimate is a linear function of the frequency of the search query
- ▶ **Hyperparameters:** $\{\alpha, \beta\}$
these are unknown and should be estimated using gradient descent
- ▶ **Loss function:** $\mathcal{J}(\alpha, \beta) = \frac{1}{2n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$
- ▶ **Goal:** $\min_{\alpha, \beta} \mathcal{J}(\alpha, \beta)$

Supervised learning — OLS with gradient descent

In our example, we are modelling a flu rate y_i using the frequency of a search query x_i

- ▶ Start with some initial values for α and β denoted by α_0 and β_0 , respectively
- ▶ In iteration $t + 1$ of the gradient descent algorithm, update α and β with:

$$\alpha_{t+1} = \alpha_t - \ell \frac{\partial \mathcal{J}(\alpha, \beta)_t}{\partial \alpha} \quad \text{and} \quad \beta_{t+1} = \beta_t - \ell \frac{\partial \mathcal{J}(\alpha, \beta)_t}{\partial \beta}$$

where ℓ often $\in (0,1)$ denotes the learning rate we want to impose

- ▶ NB: both derivatives update in iteration $t + 1$ based on values from iteration t
- ▶ Repeat until convergence

Supervised learning — OLS with gradient descent, the derivatives

Loss function: $\mathcal{J}(\alpha, \beta) = \frac{1}{2n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$ n samples, $2n$ is a convention, $\mathcal{J} = \text{MSE}/2$

$$= \frac{1}{2n} \sum_{i=1}^n (\alpha x_i + \beta - y_i)^2$$

$$\frac{\partial \mathcal{J}(\alpha, \beta)}{\partial \alpha} = \frac{1}{2n} \sum_{i=1}^n (2(\alpha x_i + \beta - y_i) x_i) = \frac{1}{n} \sum_{i=1}^n ((\alpha x_i + \beta - y_i) x_i)$$

$$\frac{\partial \mathcal{J}(\alpha, \beta)}{\partial \beta} = \frac{1}{n} \sum_{i=1}^n (\alpha x_i + \beta - y_i)$$

$$\mathcal{J}(\mathbf{w}, \beta) = \frac{1}{2n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

What if we had m predictors?


$$= \frac{1}{2n} \sum_{i=1}^n (w_1 x_{i,1} + \dots + w_m x_{i,m} + \beta - y_i)^2$$

$$\frac{\partial \mathcal{J}(\mathbf{w}, \beta)}{\partial w_j} = \frac{1}{n} \sum_{i=1}^n \left((w_1 x_{i,1} + \dots + w_m x_{i,m} + \beta - y_i) x_{i,j} \right)$$

$$\frac{\partial \mathcal{J}(\mathbf{w}, \beta)}{\partial \beta} = ?$$

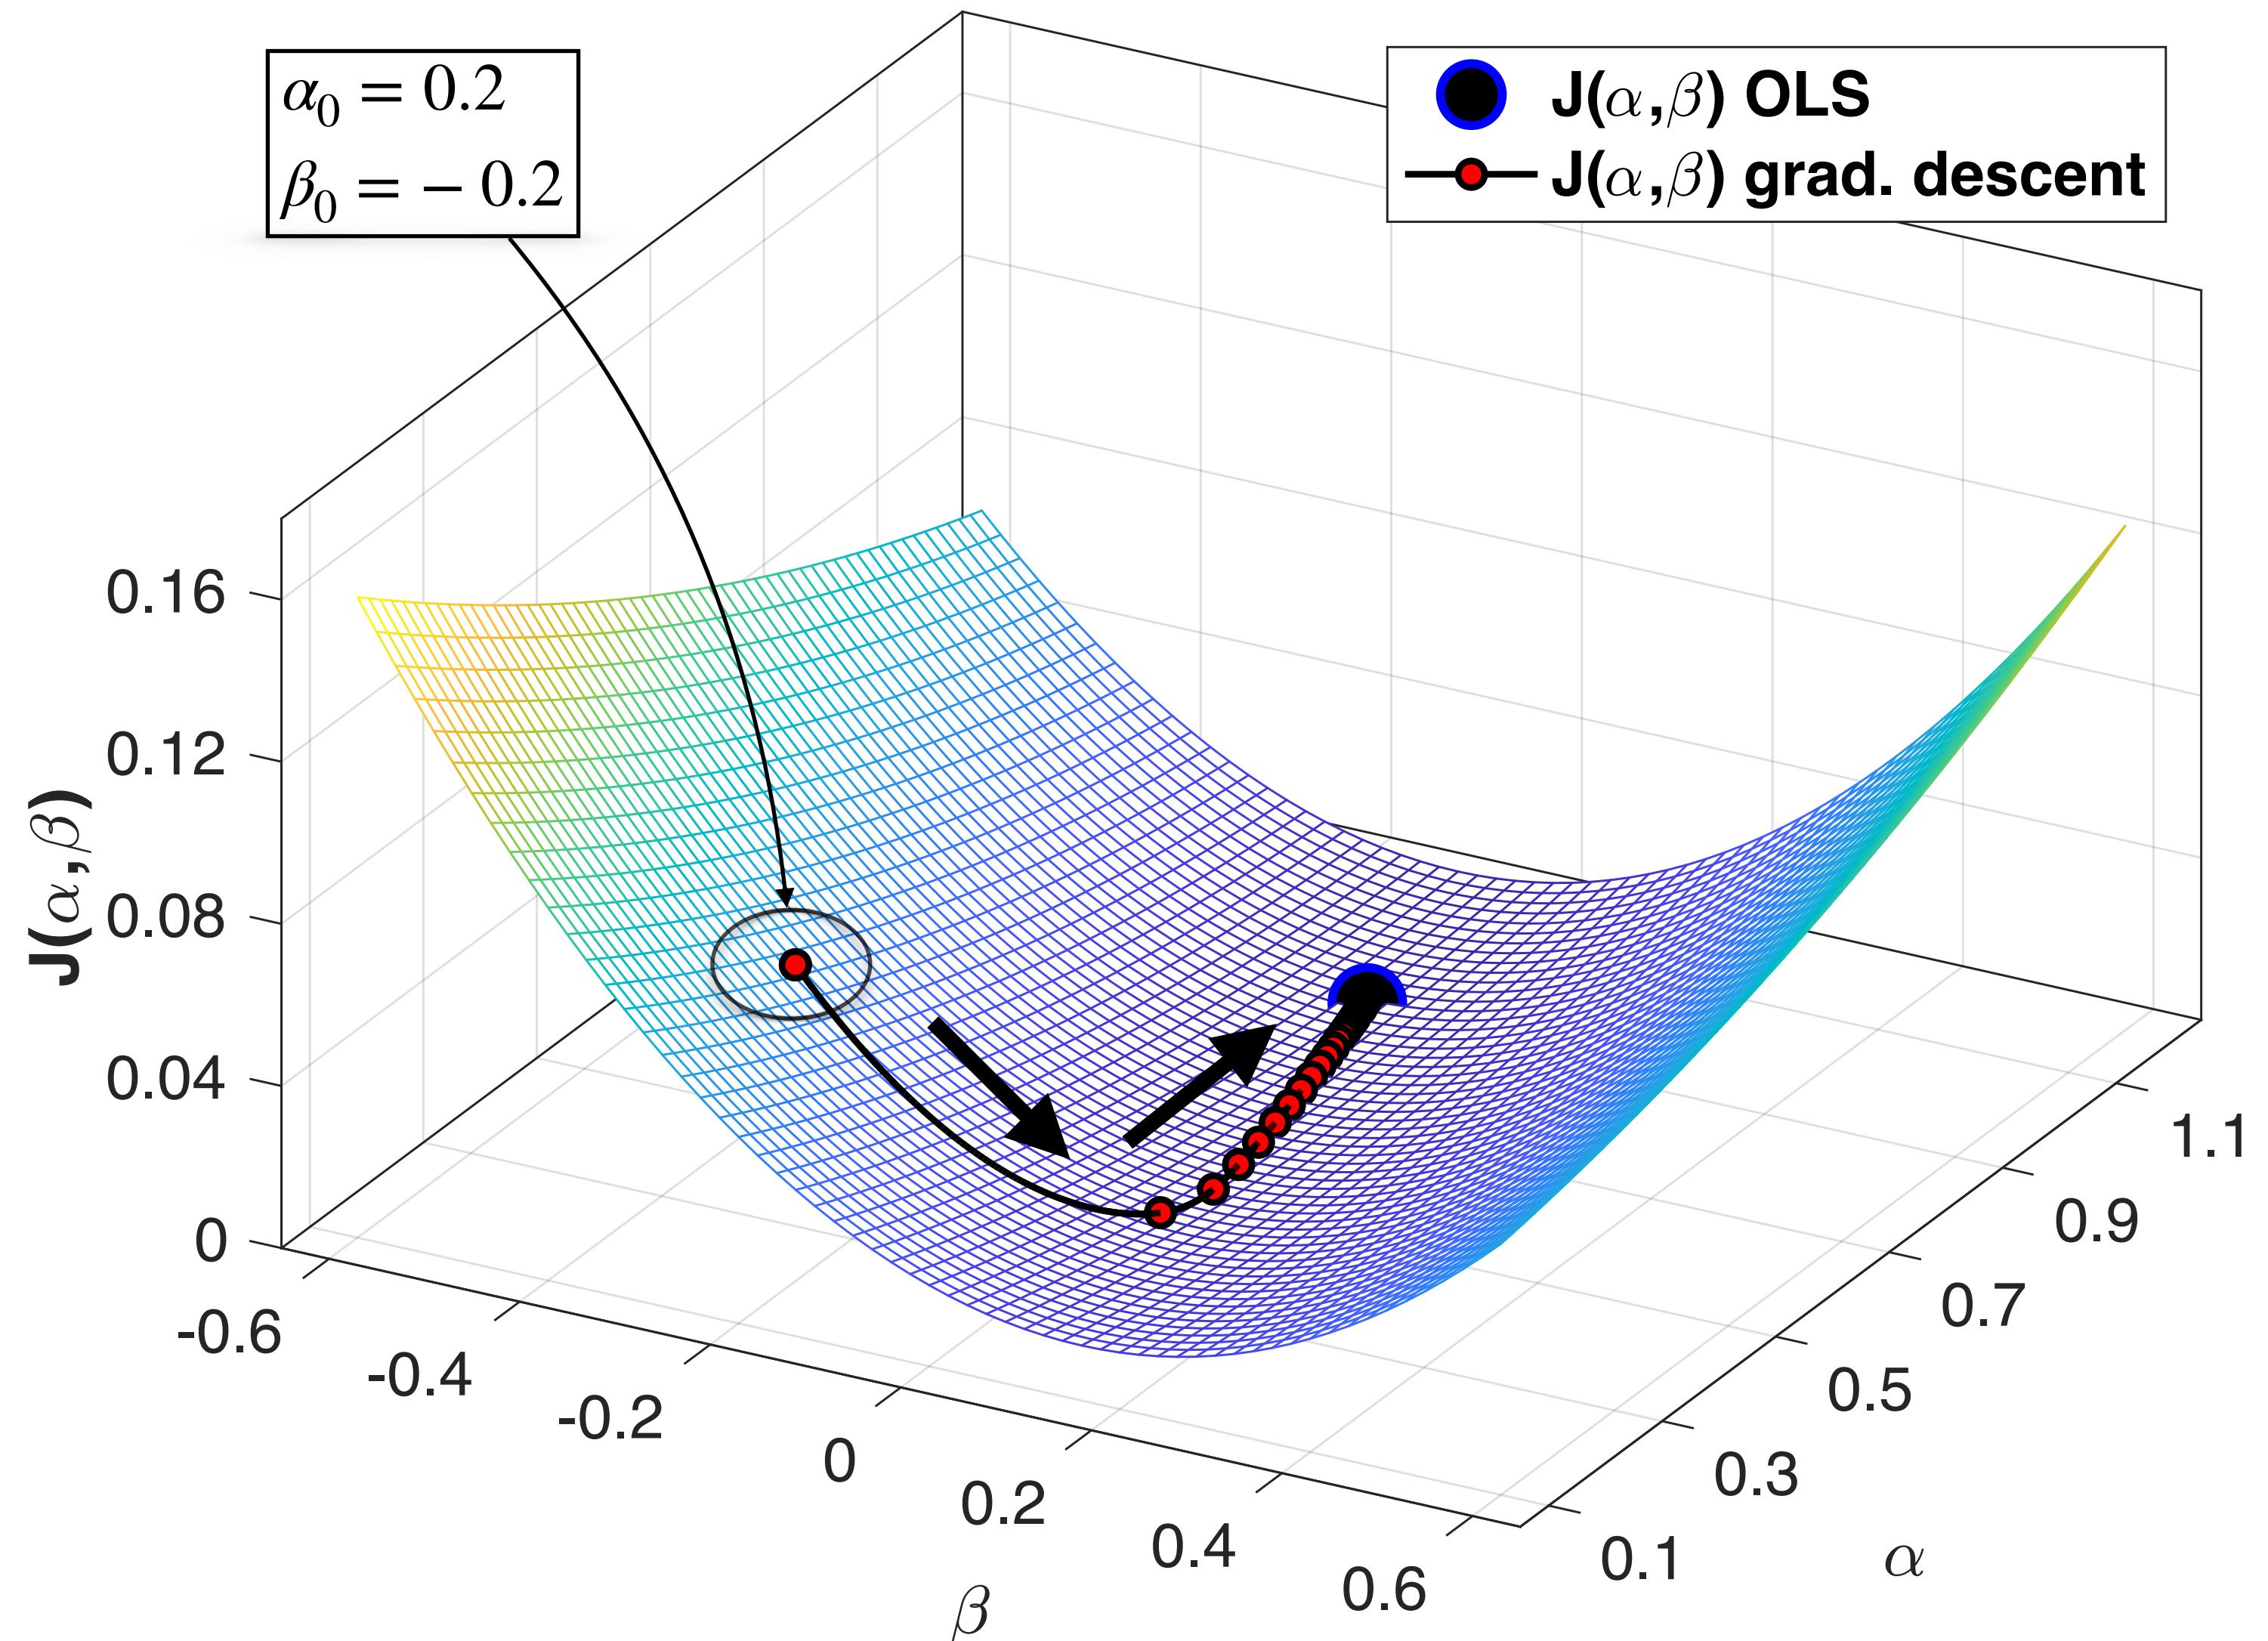
Supervised learning – OLS with gradient descent

- ▶ OLS example: inferring flu prevalence based on the frequency of 1 search query
- ▶ Let's explore the space of hyperparameter values for OLS $\{\alpha, \beta\}$ and the corresponding values of the loss function $\mathcal{J}(\alpha, \beta)$ – 3-dimensional plot (surface or mesh plot)
- ▶ Convex loss (**easier task?**)
- ▶ Big (half) dot/ball denotes the exact OLS solution (*no gradient descent used*)




Supervised learning — OLS with gradient descent

- ▶ Let's start from a point in the grid, set some initial values for the hyperparameters and attempt to solve this with coordinate descent
- ▶ $\alpha_0 = 0.2, \beta_0 = -0.2$
- ▶ $\ell = 0.02$ (learning rate)
- ▶ Convergence criterion: How much has $\mathcal{J}(\alpha, \beta)$ changed in the past k iterations?
- ▶ Gradient descent's solution almost identical to exact OLS solution (expected?)



Supervised learning – OLS with gradient descent

- ▶ Let's change the starting point
- ▶ $\alpha_0 = 1.1, \beta_0 = 0.5$
- ▶ $\ell = 0.02$ (same learning rate)
- ▶ In this case, it does not affect our solution (why?)



Supervised learning — Gradient descent, general remarks

- ▶ Effect of **learning rate** ℓ
 - if it is too small, gradient descent can be slow
 - if it is too large, gradient descent may fail to converge (overshoots the minimum)
 - adaptive learning rate (*by using line search*)
- ▶ Different initialisations might help get past local optima
- ▶ **Batch** gradient descent (*presented today*): use the entire training set for gradient updates
 - guaranteed convergence to a local minimum
 - slow on large problems (e.g. neural networks)
- ▶ **Stochastic** gradient descent: use one training sample for gradient updates
 - faster convergence on large redundant data sets
 - hard to reach high accuracy
- ▶ **Mini-batch** gradient descent: use a subset of the training set for gradient updates
 - very common in neural network training
 - better in avoiding local minima
 - what is the best mini-batch size (number of training samples to use)?

Supervised learning

► Regression

estimate / predict a continuous output / target variable

i.e. learn $f: \mathbf{X} \in \mathbb{R}^{n \times m} \rightarrow \mathbf{y} \in \mathbb{R}^n$

Examples: predict a time series trend (e.g. in finance), estimate the prevalence of a condition in epidemiology

► Classification

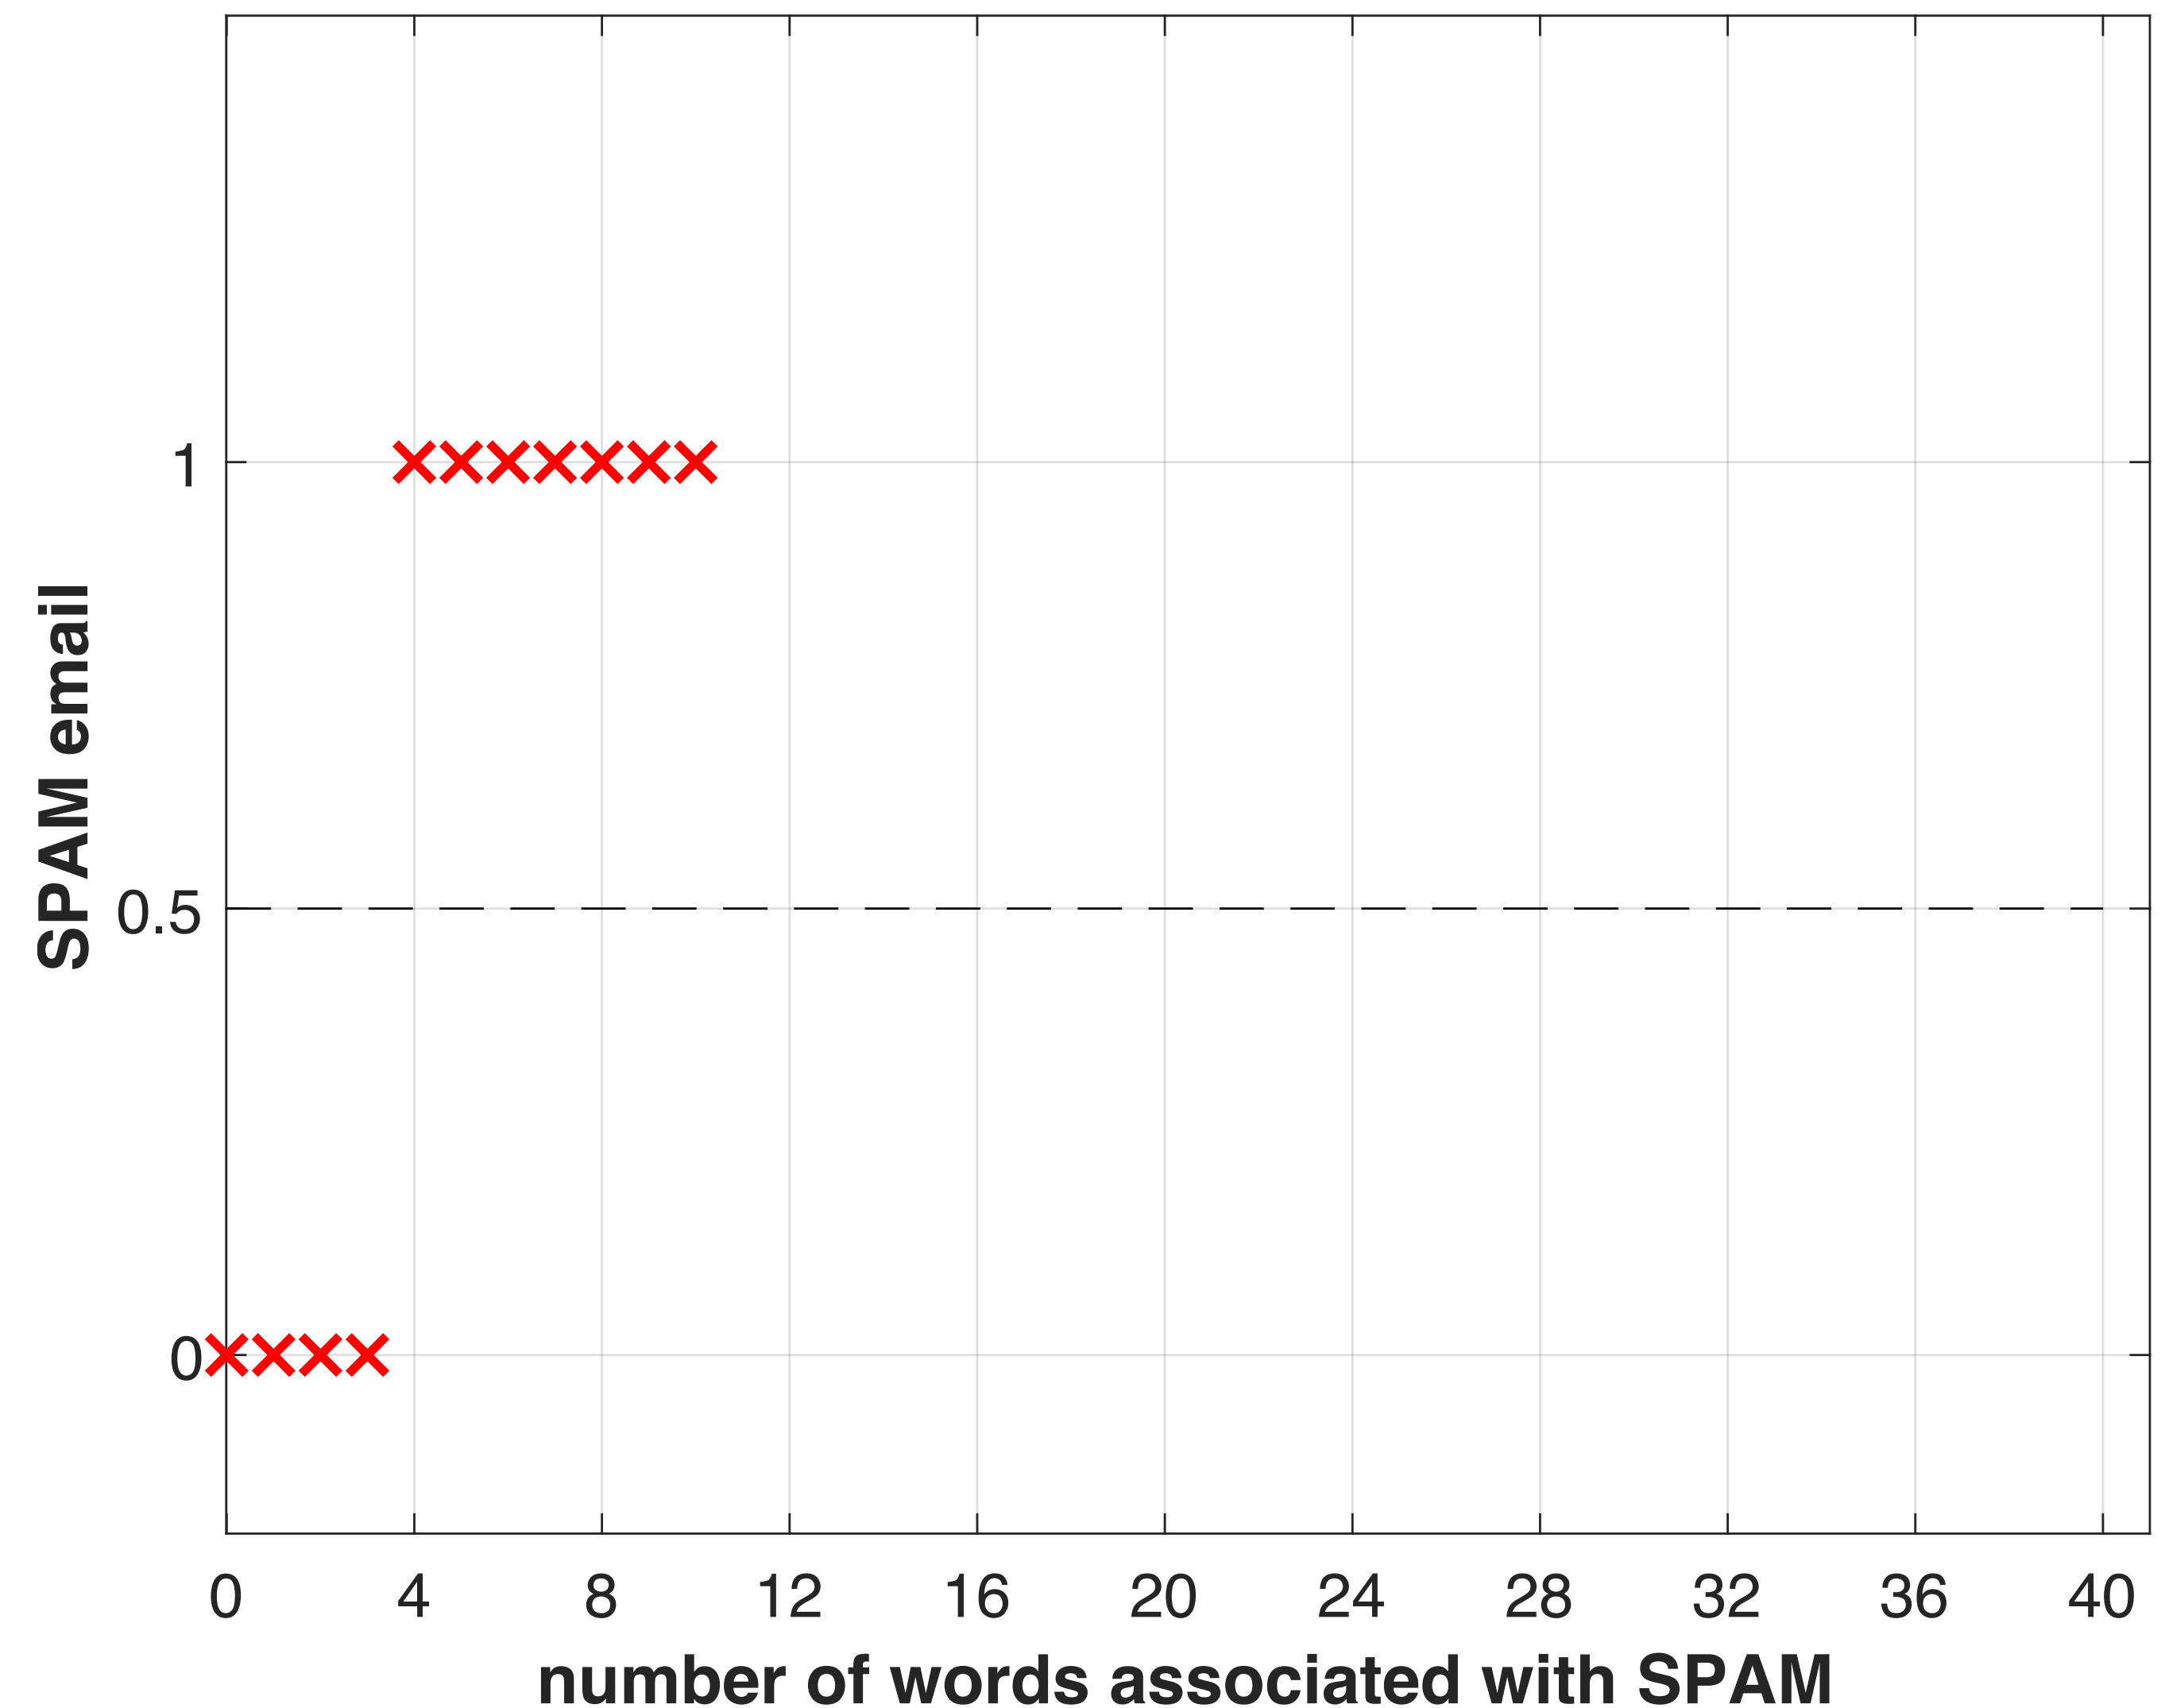
estimate a set of C unordered (and mutually exclusive) labels / classes

i.e. learn $f: \mathbf{X} \in \mathbb{R}^{n \times m} \rightarrow \mathbf{y} \in \{1, 2, \dots, C\}$

Examples: detect spam email, medical imaging, text/document classification

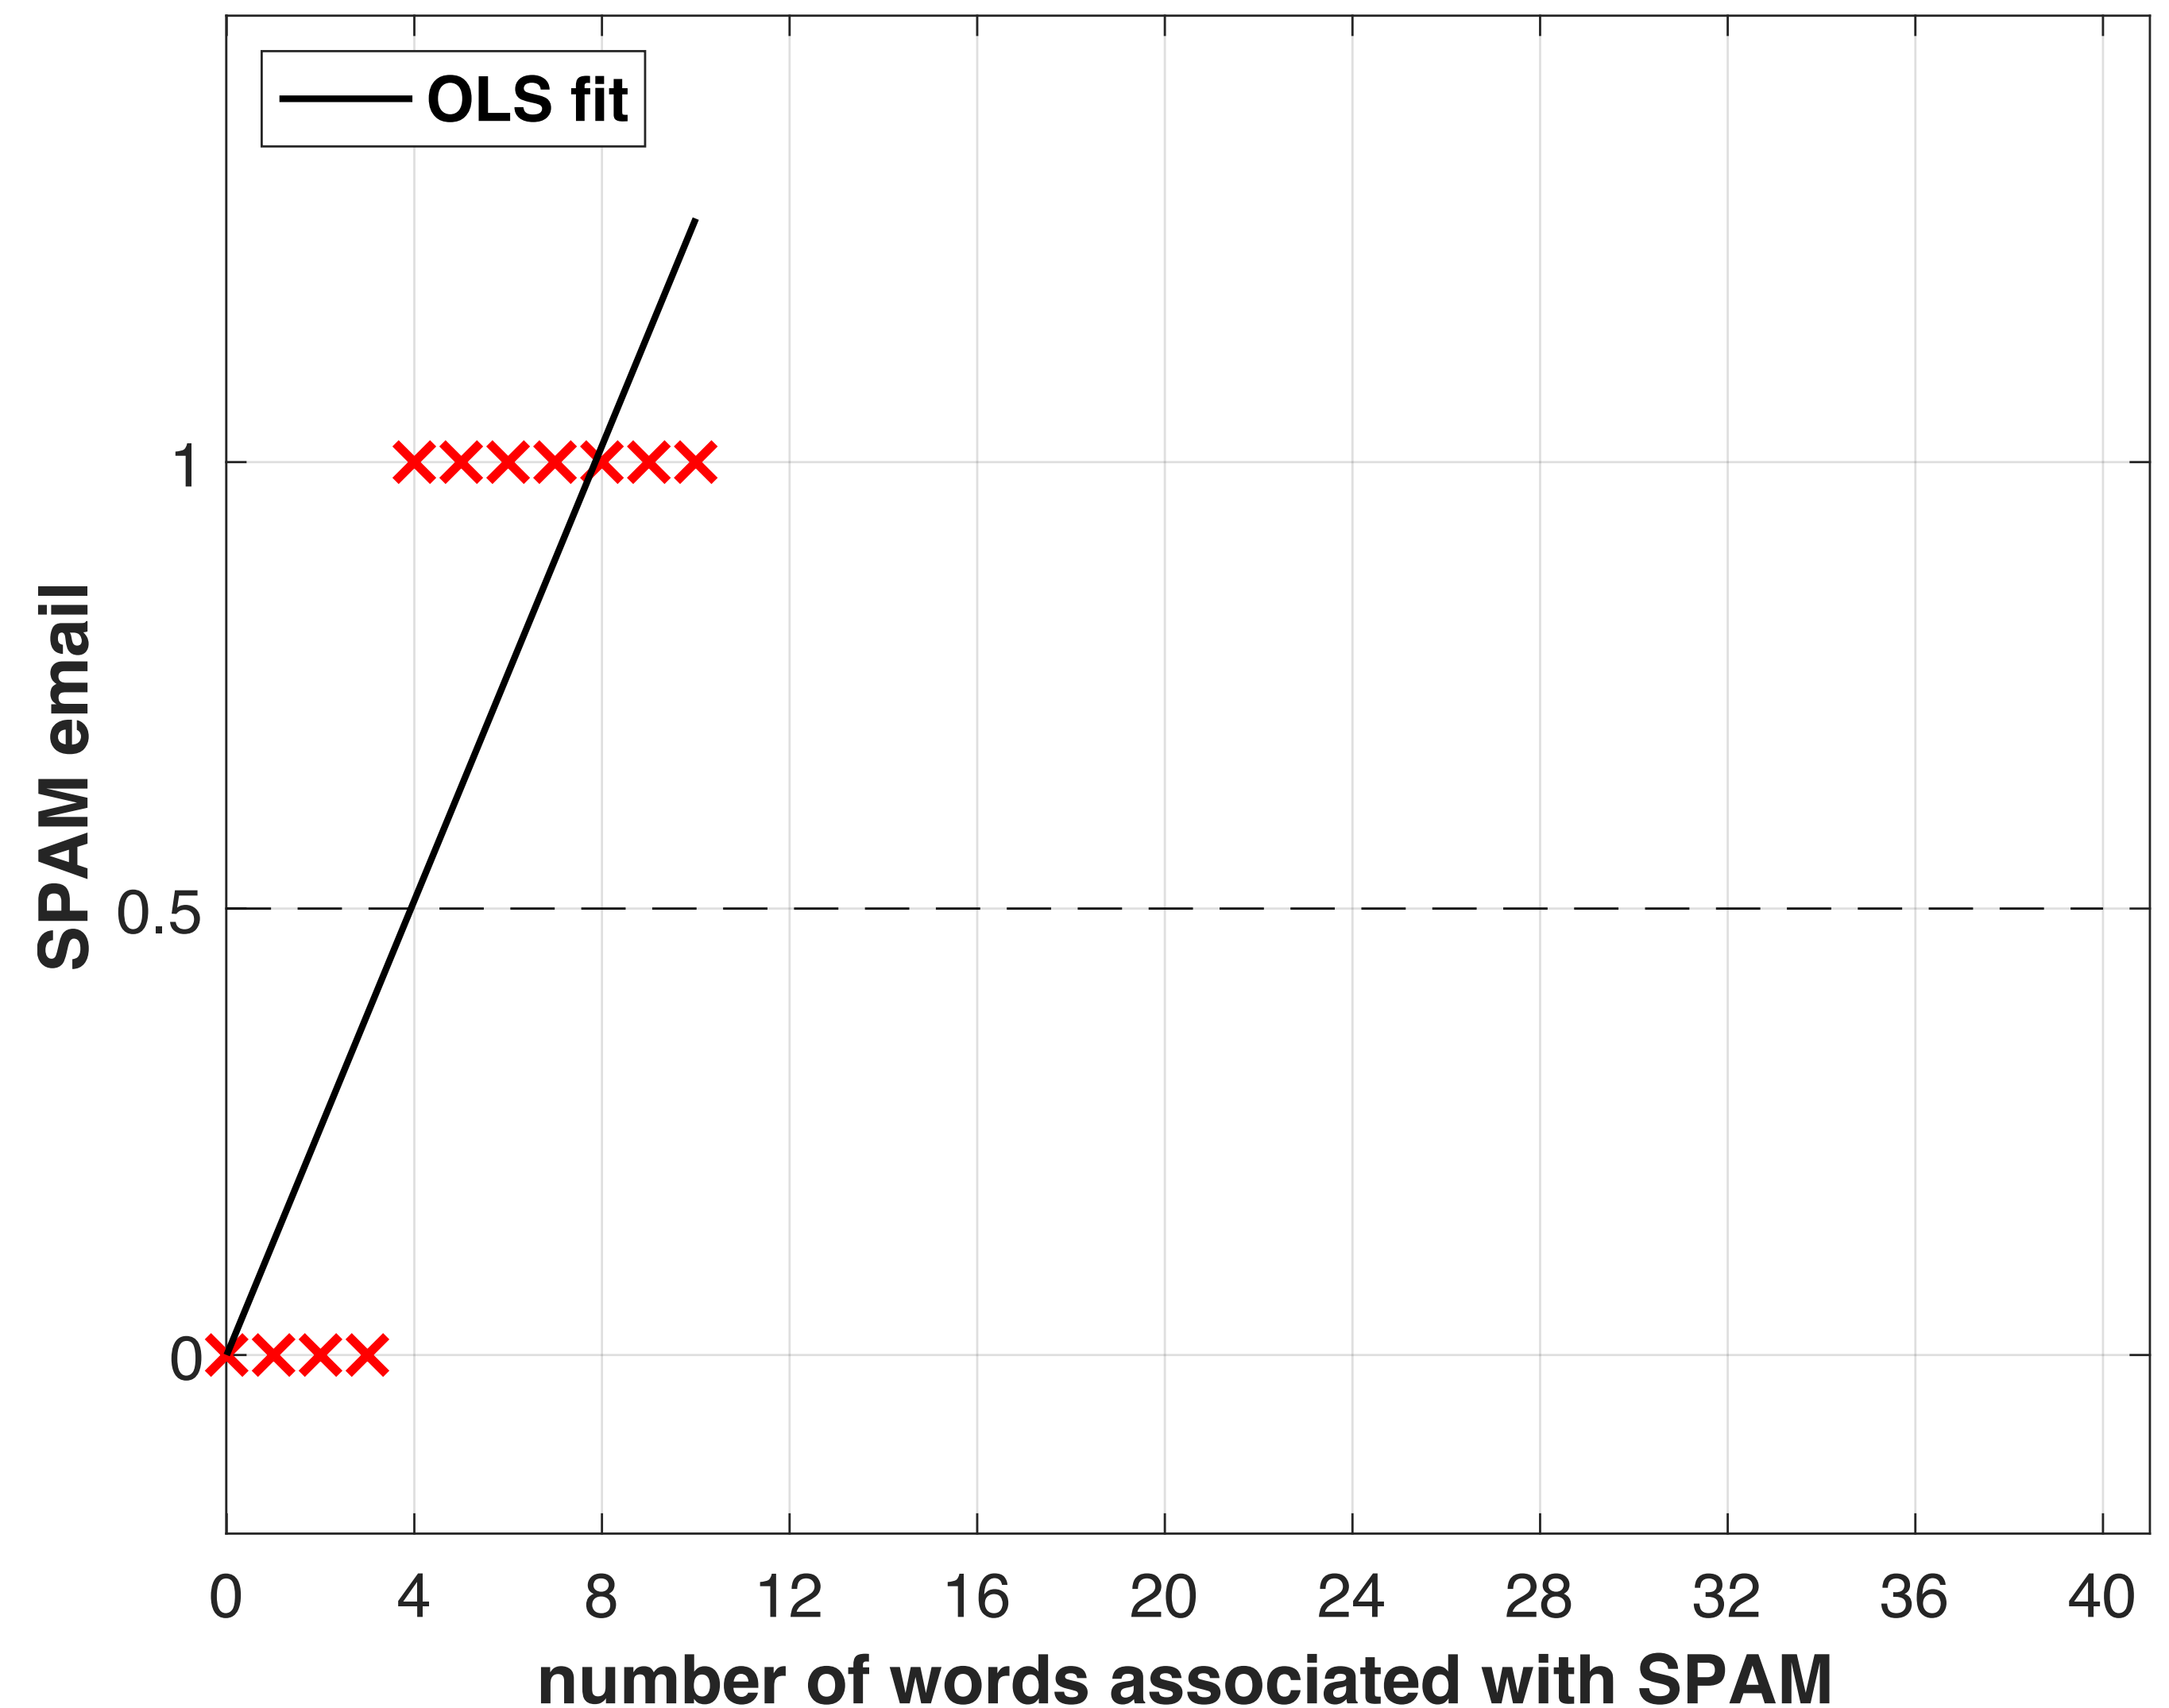
Supervised learning – Binary classification

- ▶ **Binary classification** means that we only have two label categories, e.g.
 - > *spam* vs. *not spam* email
 - > *relevant* vs. *not relevant* document
- ▶ if $f_{\mathbf{w}}(x_i) \geq 0.5$, then SPAM
if $f_{\mathbf{w}}(x_i) < 0.5$, then not SPAM
- ▶ What if we used OLS to learn $f_{\mathbf{w}}$?




Supervised learning — Binary classification

- ▶ Binary classification means that we only have two label categories, e.g.
 - > *spam* vs. *not spam* email
 - > *relevant* vs. *not relevant* document
- ▶ if $f_{\mathbf{w}}(x_i) \geq 0.5$, then SPAM
if $f_{\mathbf{w}}(x_i) < 0.5$, then not SPAM
- ▶ What if we used OLS to learn $f_{\mathbf{w}}$?
Looks perfectly fine?



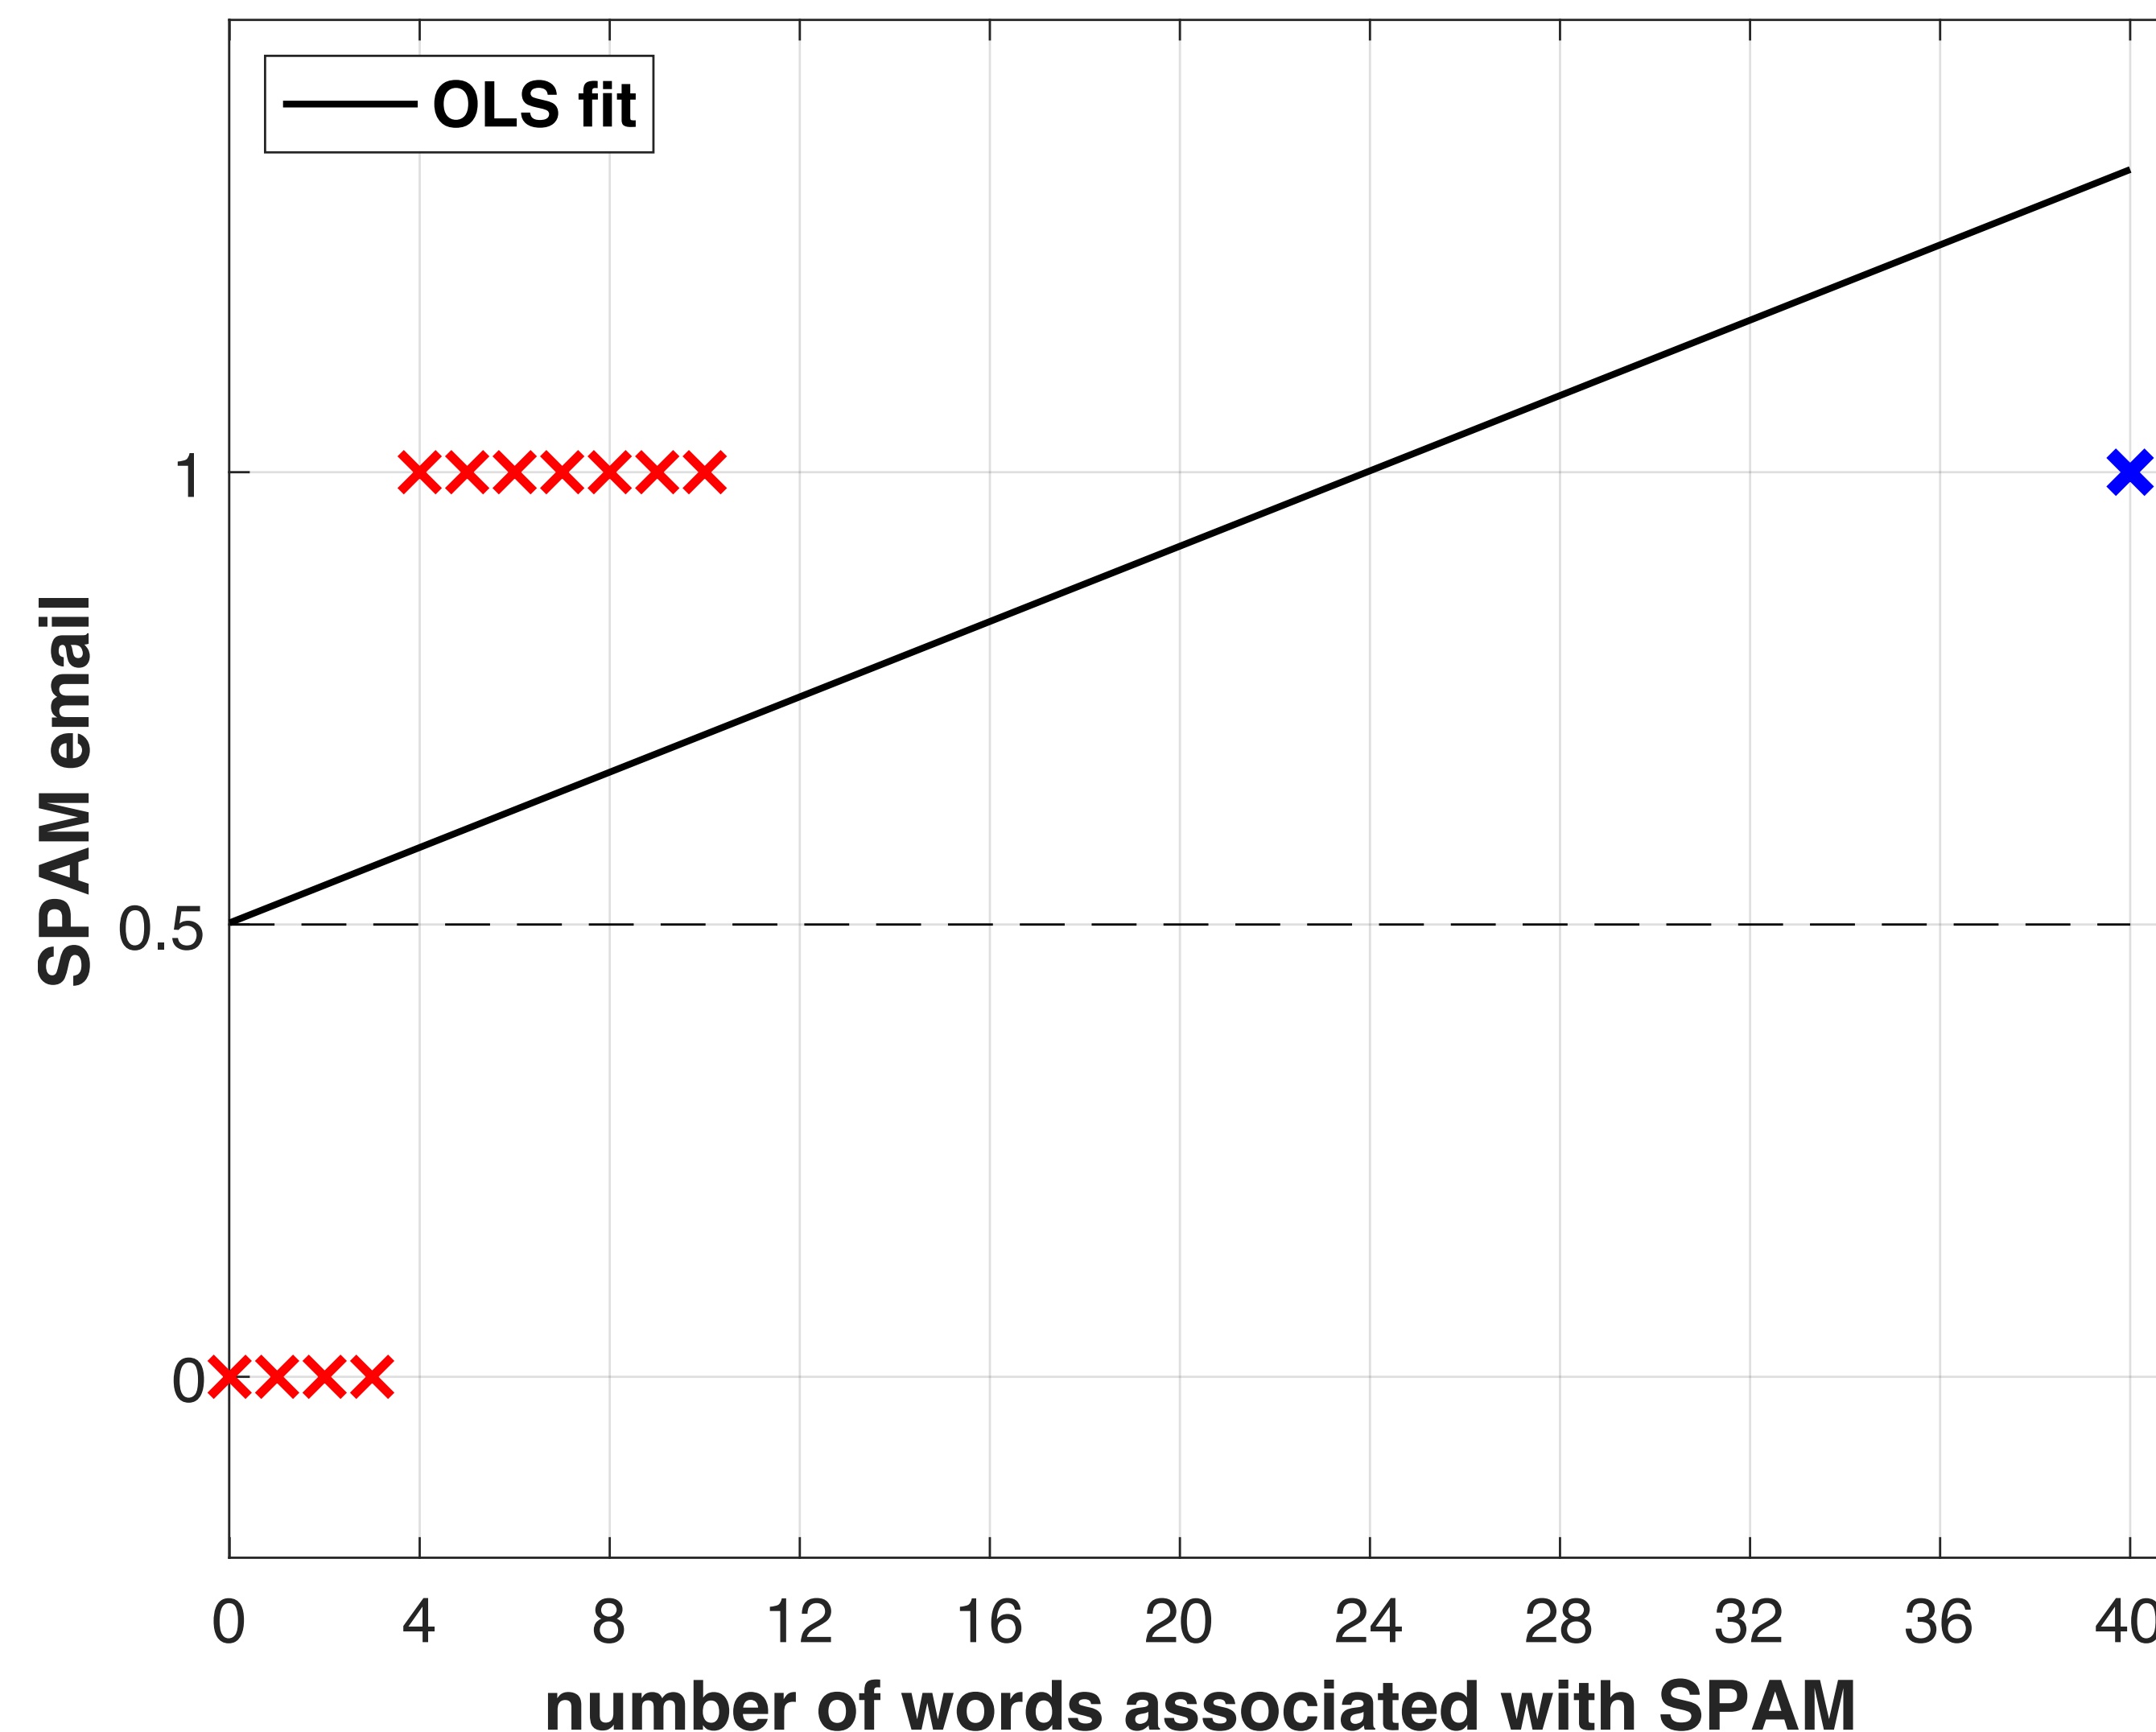
Supervised learning – Binary classification

- ▶ **Binary classification** means that we only have two label categories, e.g.
 - > *spam* vs. *not spam* email
 - > *relevant* vs. *not relevant* document
- ▶ if $f_{\mathbf{w}}(x_i) \geq 0.5$, then SPAM
if $f_{\mathbf{w}}(x_i) < 0.5$, then not SPAM
- ▶ What if we used OLS to learn $f_{\mathbf{w}}$?
Looks perfectly fine?
- ▶ Let's add one more observation to our data. How would that affect our OLS classifier?




Supervised learning — Binary classification

- ▶ **Binary classification** means that we only have two label categories, e.g.
 - > *spam* vs. *not spam* email
 - > *relevant* vs. *not relevant* document
- ▶ if $f_{\mathbf{w}}(x_i) \geq 0.5$, then SPAM
if $f_{\mathbf{w}}(x_i) < 0.5$, then not SPAM
- ▶ What if we used OLS to learn $f_{\mathbf{w}}$?
Looks perfectly fine?
- ▶ Let's add one more observation to our data. How would that affect our OLS classifier? **Not great!**




Supervised learning — Binary classification

- ▶ **Binary classification** means that we only have two label categories, e.g.
 - > *spam* vs. *not spam* email
 - > *relevant* vs. *not relevant* document
- ▶ if $f_w(x_i) \geq 0.5$, then SPAM
if $f_w(x_i) < 0.5$, then not SPAM
- ▶ What if we used OLS to learn f_w ?
Looks perfectly fine?
- ▶ Let's add one more observation to our data. How would that affect our OLS classifier? **Not great!**
- ▶ It is not impossible to separate these classes — we just need a **different function**.




- ▶ Assume we have an m -dimensional observation $\mathbf{x} \in \mathbb{R}^m$
- ▶ We want $0 \leq f_{\mathbf{w}}(\mathbf{x}) \leq 1$, where $\mathbf{w} \in \mathbb{R}^m$ are the corresponding weights
- ▶ Sigmoid or logistic function
$$\sigma(z) = \frac{1}{1 + e^{-z}}$$
- ▶ $f_{\mathbf{w}}(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x}) \in [0,1]$
it can be seen as a probability

Sigmoid / logistic function




Supervised learning — Logistic regression's loss function

- ▶ Cross-entropy loss function between the projection of \mathbf{x} and label $y \in \{0,1\}$
- ▶ $\mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) = -\ln(\sigma_{\mathbf{w}}(\mathbf{x}))$ if $y = 1$
 $\mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) = -\ln(1 - \sigma_{\mathbf{w}}(\mathbf{x}))$ if $y = 0$
- ▶ Derivation from Bernoulli distribution
- ▶ Intuitively
 - we want a loss that is easy to differentiate
 - if $y = 1, \sigma_{\mathbf{w}}(\mathbf{x}) \rightarrow 1: \mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) \rightarrow 0$
 - if $y = 1, \sigma_{\mathbf{w}}(\mathbf{x}) \rightarrow 0: \mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) \rightarrow \infty$



Supervised learning — Logistic regression's loss function

- ▶ Cross-entropy loss function between the projection of \mathbf{x} and label $y \in \{0,1\}$
- ▶ $\mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) = -\ln(\sigma_{\mathbf{w}}(\mathbf{x}))$ if $y = 1$
 $\mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) = -\ln(1 - \sigma_{\mathbf{w}}(\mathbf{x}))$ if $y = 0$
- ▶ Derivation from Bernoulli distribution
- ▶ Intuitively
 - we want a loss that is easy to differentiate
 - if $y = 0, \sigma_{\mathbf{w}}(\mathbf{x}) \rightarrow 0: \mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) \rightarrow 0$
 - if $y = 0, \sigma_{\mathbf{w}}(\mathbf{x}) \rightarrow 1: \mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) \rightarrow \infty$



Cross-entropy loss function

$$\begin{aligned}\mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) &= -\ln(\sigma_{\mathbf{w}}(\mathbf{x})) && \text{if } y = 1 \\ \mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) &= -\ln(1 - \sigma_{\mathbf{w}}(\mathbf{x})) && \text{if } y = 0\end{aligned}$$

Logistic (sigmoid) function

$$\sigma_{\mathbf{w}}(\mathbf{x}_i) = \left(1 + e^{-\mathbf{w}^\top \mathbf{x}_i}\right)^{-1}$$

Our label y_i is either 1 or 0 for all our observations. So, for each observation only one part of the loss function is activated / used. Since we have n observations the loss function takes the form:

$$\mathcal{J}(\mathbf{w}) = -\frac{1}{n} \sum_{i=1}^n \left[y_i \ln \sigma_{\mathbf{w}}(\mathbf{x}_i) + (1 - y_i) \ln (1 - \sigma_{\mathbf{w}}(\mathbf{x}_i)) \right]$$

Supervised learning — Logistic regression with gradient descent

Cross-entropy loss function

$$\begin{aligned}\mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) &= -\ln(\sigma_{\mathbf{w}}(\mathbf{x})) && \text{if } y = 1 \\ \mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) &= -\ln(1 - \sigma_{\mathbf{w}}(\mathbf{x})) && \text{if } y = 0\end{aligned}$$

Logistic (sigmoid) function

$$\sigma_{\mathbf{w}}(\mathbf{x}_i) = \left(1 + e^{-\mathbf{w}^\top \mathbf{x}_i}\right)^{-1}$$

Combined loss function

$$\mathcal{J}(\mathbf{w}) = -\frac{1}{n} \sum_{i=1}^n \left[y_i \ln \sigma_{\mathbf{w}}(\mathbf{x}_i) + (1 - y_i) \ln (1 - \sigma_{\mathbf{w}}(\mathbf{x}_i)) \right]$$

Let's incorporate the actual value of the sigmoid function and attempt to simplify:

$$\ln(\sigma_{\mathbf{w}}(\mathbf{x}_i)) = \ln(1) - \ln(1 + e^{-\mathbf{w}^\top \mathbf{x}_i}) = -\ln(1 + e^{-\mathbf{w}^\top \mathbf{x}_i})$$

The loss function becomes:

$$\mathcal{J}(\mathbf{w}) = -\frac{1}{n} \sum_{i=1}^n \left[\mathbf{w}^\top \mathbf{x}_i y_i - \mathbf{w}^\top \mathbf{x}_i - \ln(1 + e^{-\mathbf{w}^\top \mathbf{x}_i}) \right]$$

Supervised learning — Logistic regression with gradient descent

Cross-entropy loss function

$$\begin{aligned}\mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) &= -\ln(\sigma_{\mathbf{w}}(\mathbf{x})) && \text{if } y = 1 \\ \mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) &= -\ln(1 - \sigma_{\mathbf{w}}(\mathbf{x})) && \text{if } y = 0\end{aligned}$$

Logistic (sigmoid) function

$$\sigma_{\mathbf{w}}(\mathbf{x}_i) = \left(1 + e^{-\mathbf{w}^\top \mathbf{x}_i}\right)^{-1}$$

Combined loss function

$$\mathcal{J}(\mathbf{w}) = -\frac{1}{n} \sum_{i=1}^n \left[\mathbf{w}^\top \mathbf{x}_i y_i - \mathbf{w}^\top \mathbf{x}_i - \ln \left(1 + e^{-\mathbf{w}^\top \mathbf{x}_i} \right) \right]$$

Partial derivative

$$\begin{aligned}\frac{\partial \mathcal{J}(\mathbf{w})}{\partial \mathbf{w}_j} &= -\frac{1}{n} \sum_{i=1}^n \left[y_i x_{i,j} - x_{i,j} + e^{-\mathbf{w}^\top \mathbf{x}_i} \left(1 + e^{-\mathbf{w}^\top \mathbf{x}_i} \right)^{-1} x_{i,j} \right] \\ &= -\frac{1}{n} \sum_{i=1}^n \left[x_{i,j} (y_i - \sigma_{\mathbf{w}}(\mathbf{x}_i)) \right]\end{aligned}$$

Supervised learning — Logistic regression with gradient descent

Cross-entropy loss function

$$\begin{aligned}\mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) &= -\ln(\sigma_{\mathbf{w}}(\mathbf{x})) && \text{if } y = 1 \\ \mathcal{L}(\sigma_{\mathbf{w}}(\mathbf{x}), y) &= -\ln(1 - \sigma_{\mathbf{w}}(\mathbf{x})) && \text{if } y = 0\end{aligned}$$

Logistic (sigmoid) function

$$\sigma_{\mathbf{w}}(\mathbf{x}_i) = \left(1 + e^{-\mathbf{w}^\top \mathbf{x}_i}\right)^{-1}$$

Combined loss function

$$\mathcal{J}(\mathbf{w})$$

Partial derivative

$$\frac{\partial \mathcal{J}(\mathbf{w})}{\partial w_j}$$

The rest is identical to the least squares example, i.e. initialise \mathbf{w} , then compute the partial derivatives for each w_j , then update w_j 's using a learning rate, and repeat until convergence.

$$\left[\frac{1}{1 + e^{-\mathbf{w}^\top \mathbf{x}_i}} - \frac{e^{-\mathbf{w}^\top \mathbf{x}_i}}{(1 + e^{-\mathbf{w}^\top \mathbf{x}_i})^2} x_{i,j} \right]$$

$$= -\frac{1}{n} \sum_{i=1}^n \left[x_{i,j} (y_i - \sigma_{\mathbf{w}}(\mathbf{x}_i)) \right]$$

- ▶ Going back to the application of estimating flu prevalence using web search activity
- ▶ Now, we want to use the frequency of 4 search queries to predict whether the flu rate in a population is above a low-epidemic threshold or not
 - binary classification task
 - $y_i = 1$, if the flu rate is above a low-epidemic threshold
 - $y_i = 0$, if the flu rate is below or equal to a low-epidemic threshold
- ▶ We have in total 104 weekly observations
 - observation matrix $\mathbf{X} \in \mathbb{R}^{104 \times 4}$
 - queries: “*how long does flu last*”, “*flu symptoms*”, “*cough flu*”, “*flu recovery*”
 - labels $\mathbf{y} \in \{0,1\}^{104}$

- ▶ We have in total 104 weekly observations
 - observation matrix $\mathbf{X} \in \mathbb{R}^{104 \times 4}$
 - queries: “*how long does flu last*”, “*flu symptoms*”, “*cough flu*”, “*flu recovery*”
 - labels $\mathbf{y} \in \{0,1\}^{104}$
- ▶ 4-fold cross validation
 - form 4 folds (equally sized baskets) of the data,
 - train a classifier using 3 of them, test (evaluate) on the remaining 1
 - report average performance metrics

Binary classification — Basic performance metrics

► $\hat{\mathbf{y}} \in \{0,1\}^n$ denotes our predictions and $\mathbf{y} \in \{0,1\}^n$ the correct labels

► $\text{accuracy} = \frac{\text{number of times } \hat{y}_i = y_i}{n}$

► $\text{precision} = \frac{\text{number of times } \hat{y}_i = 1 \text{ AND } \hat{y}_i = y_i}{\text{number of times } \hat{y}_i = 1}$

► $\text{recall} = \frac{\text{number of times } \hat{y}_i = 1 \text{ AND } \hat{y}_i = y_i}{\text{number of times } y_i = 1}$

► F_1 score is the harmonic mean between precision and recall

$$F_1 \text{ score} = 2 \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

Supervised learning — Logistic regression, example

logistic regression performance metrics

accuracy = 0.923 (0.070)

precision = 0.902 (0.121)

recall = 0.844 (0.120)

F_1 score = 0.871(0.116)

logistic regression weights

flu symptoms: 78.058

how long does flu last: 24.537

flu recovery: 3.8977

cough flu: – 14.663

Download the data from
dropbox.com/s/rgyg190whw26qrj/data-COMP0084-intro-to-ml.zip?dl=0
and try it at home!

Multi-class classification

- ▶ Binary classification is the simplest classification case — we often have more than two labels, i.e. most tasks require multi-class classification
- ▶ We can use different classifiers (machine learning models) that support multi-class classification such as neural network architectures and generative models
- ▶ We can also use a binary classifier
 - **one vs. rest** strategy: n classes require n classifiers to be trained
highest score determines the classification label
 - **one vs. one** strategy: n classes require $\frac{n(n-1)}{2}$ classifiers to be trained
voting scheme, class with the most votes wins


- ▶ In the previous machine learning paradigms we had an input \mathbf{X} and an output \mathbf{y} and we wanted to learn $f: \mathbf{X} \rightarrow \mathbf{y}$
- ▶ In unsupervised learning, there are no particular outputs or, better, response variables that we can associate our inputs with
- ▶ Our goal now is different: we want to extract some kind of pattern (a rule, an intrinsic structure) from a data set (a set of observations \mathbf{X})

Unsupervised learning

- ▶ Is association rule mining a form of unsupervised learning?
- ▶ Dimensionality reduction, principal component analysis
- ▶ Clustering: almost synonymous with unsupervised learning
- ▶ Clustering groups similar observations (features) together into **clusters**

Unsupervised learning – Clustering

- ▶ A **cluster** contains data instances that are similar to each other (or if you visualise this are very close to each other in a vector space) – in very lay terms, different clusters are supposed to be capturing a different part of this vector space
- ▶ So, clustering is a grouping of data objects such that the objects within a group are similar (or related) to one another and different from (or unrelated to) the objects in other groups
- ▶ The plot shows 3 very visible clusters



Clustering — Where is it being used?


- ▶ Computational biology, e.g. understand properties of genes
- ▶ Medicine, e.g. in medical imaging
- ▶ Marketing, e.g. segment customers according to their underlying characteristics, then conduct targeted marketing
- ▶ Document clustering, topic models, text clustering in general
- ▶ Applicable to tasks that require “pattern analysis” and in many different research disciplines for analysing outcomes (e.g. in psychology, sociology, computer science, neuroscience)

Clustering – Some further key aspects

- ▶ Many different clustering algorithms / methods
 - partitional, hierarchical, hard/soft, generative, and even supervised
- ▶ A distance (dissimilarity) or a similarity function is often a key component for determining clusters
- ▶ Clustering goal is to maximise the distance between different clusters (inter-cluster distance) and at the same time to minimise the distance of elements in a cluster (intra-cluster distance)
- ▶ The quality of a clustering outcome depends on the algorithm, the distance function, and eventually the specifics of an application
- ▶ However, determining the actual quality of a cluster is not always an easy task given the lack of supervision

Clustering – Outliers

- ▶ Outliers are objects that do not belong to any cluster or form clusters of very small cardinality
- ▶ In some applications we are actually interested in discovering outliers, not clusters



Clustering — Distance / similarity functions

- ▶ Let's assume we want to compare two n -dimensional observations, \mathbf{x} and \mathbf{z}
- ▶ Let's also assume first that both \mathbf{x} and \mathbf{z} contain discrete values; these can be binary values (0 or 1), or specific element identifiers
- ▶ The Jaccard similarity provides a simple way to compare these observations

$$J_{\text{sim}}(\mathbf{x}, \mathbf{z}) = \frac{|\mathbf{x} \cap \mathbf{z}|}{|\mathbf{x} \cup \mathbf{z}|}$$

- ▶ Comparisons are element-wise
- ▶ The Jaccard similarity can (by definition) take values from 0 to 1
- ▶ It can be turned to a distance metric: $J_{\text{dist}}(\mathbf{x}, \mathbf{z}) = 1 - J_{\text{sim}}(\mathbf{x}, \mathbf{z})$
- ▶ if $\mathbf{x} = \begin{bmatrix} 1 & 0 & 0 & 1 & 1 & 1 \end{bmatrix}$
 $\mathbf{z} = \begin{bmatrix} 0 & 1 & 1 & 0 & 1 & 0 \end{bmatrix}$ then $J_{\text{sim}}(\mathbf{x}, \mathbf{z}) = 1/6$ and $J_{\text{dist}}(\mathbf{x}, \mathbf{z}) = 5/6$

- ▶ Let's assume we want to compare two n -dimensional observations, \mathbf{x} and \mathbf{z}
- ▶ Let's now assume that both \mathbf{x} and $\mathbf{z} \in \mathbb{R}^n$

- ▶ Recall the L_p -norm definition: $\|\mathbf{x}\|_p = \left(\sum_{i=1}^n |x_i|^p \right)^{\frac{1}{p}}$

- ▶ Popular distance measures stem from this – the input now is the difference of the vectors we want to compare – this is also known as the **Minkowski distance**

$$L_p(\mathbf{x}, \mathbf{z}) = \left(|x_1 - z_1|^p + |x_2 - z_2|^p + \dots + |x_n - z_n|^p \right)^{1/p} = \|\mathbf{x} - \mathbf{z}\|_p$$

- ▶ Let's assume we want to compare two n -dimensional observations, \mathbf{x} and $\mathbf{z} \in \mathbb{R}^n$

- ▶ **Minkowski distance**

$$L_p(\mathbf{x}, \mathbf{z}) = \left(|x_1 - z_1|^p + |x_2 - z_2|^p + \dots + |x_n - z_n|^p \right)^{1/p} = \|\mathbf{x} - \mathbf{z}\|_p$$

- ▶ For different values of $p \in \mathbb{N}_{>0}$ we can obtain common distance functions

- ▶ $p = 1$, **Manhattan or city block distance** or L_1 -norm

$$L_1(\mathbf{x}, \mathbf{z}) = |x_1 - z_1| + |x_2 - z_2| + \dots + |x_n - z_n|$$

- ▶ $p = 2$, **Euclidean distance** or L_2 -norm

$$L_2(\mathbf{x}, \mathbf{z}) = \left[(x_1 - z_1)^2 + (x_2 - z_2)^2 + \dots + (x_n - z_n)^2 \right]^{1/2} = \sqrt{(x_1 - z_1)^2 + (x_2 - z_2)^2 + \dots + (x_n - z_n)^2}$$

- ▶ Let's assume we want to compare two n -dimensional observations, \mathbf{x} and $\mathbf{z} \in \mathbb{R}^n$

- ▶ **Minkowski distance**

$$L_p(\mathbf{x}, \mathbf{z}) = \left(|x_1 - z_1|^p + |x_2 - z_2|^p + \dots + |x_n - z_n|^p \right)^{1/p} = \|\mathbf{x} - \mathbf{z}\|_p$$

- ▶ For different values of $p \in \mathbb{N}_{>0}$ we can obtain common distance functions
- ▶ We can also weight distances, if we want to give more importance to certain features, e.g.

$$L_2(\mathbf{x}, \mathbf{z}) = \sqrt{w_1 (x_1 - z_1)^2 + w_2 (x_2 - z_2)^2 + \dots + w_n (x_n - z_n)^2}$$

- ▶ Different approaches to clustering
 - relation between objects and classes (exclusive vs. overlapping)
 - relation between classes and classes (ordered vs. flat)
- ▶ Today we are going to see k -means: driven by the relationship to cluster representatives (or means), partitional clustering algorithm
- ▶ Construct a partition of a set of n features (objects) into a set of k clusters
 - each object belongs to exactly one cluster (hard clustering)
 - the number of clusters (k) is a setting given in advance

Clustering algorithms – k -means

- ▶ Let's assume we have a set of n m -dimensional observations, i.e. a matrix $\mathbf{X} \in \mathbb{R}^{n \times m}$
 - the number of dimensions = number of features (m)
 - a feature i is represented by the i -th column of \mathbf{X} , the n -dimensional vector $\mathbf{x}_{:,i}$
 - we want to partition the features of \mathbf{X} into k clusters
- 1. Randomly select k data points (seeds) to be the initial cluster centres. We call these centres **centroids** and in practice they are n -dimensional vectors (same size as the number of observations). Centroid j is denoted by $\mathbf{c}_j \in \mathbb{R}^n$
- 2. Assign each feature $\mathbf{x}_{:,i}$ to its closest centroid \mathbf{c}_j
- 3. Re-compute centroids by averaging across their members
- 4. If a convergence criterion is not met, go back to step 2

Clustering algorithms – k -means convergence criteria

- ▶ no or minimum re-assignments of data points to different clusters
- ▶ no or minimum change of centroids
- ▶ minimum decrease in the following cost function – the distance of all features from their centroids has converged to a minimum (C_j denotes cluster j)

$$\sum_{j=1}^k \sum_{\mathbf{x}_{:,i} \in C_j} \text{dist}(\mathbf{x}_{:,i}, \mathbf{c}_j)$$

- ▶ we can use different distance functions, the most common being the Euclidean distance squared, i.e.

$$\sum_{j=1}^k \sum_{\mathbf{x}_{:,i} \in C_j} \|\mathbf{x}_{:,i} - \mathbf{c}_j\|_2^2$$

Clustering algorithms – k -means


- ▶ Strengths
 - simple implementation
 - efficient, time complexity $\mathcal{O}(tkn)$, t number of iterations, k clusters, n observations
 - finds a local optimum
 - *no definitive evidence that any cluster algorithm performs better* (hard to evaluate anyway)
- ▶ Weaknesses
 - we need to specify k (the number of clusters)
 - sensitive to outliers
 - sensitive to initialisation
- ▶ Workarounds / improvements
 - multiple runs with different initialisations
 - non random initialisation, centroids set to the most distant observations (k -means++)

Clustering algorithms – k -means, an example

- ▶ Back to our web search activity data set
- ▶ 150 web search queries that are used to model flu rates in England
- ▶ Weekly frequency for 674 weeks, i.e. $\mathbf{X} \in \mathbb{R}_{\geq 0}^{674 \times 150}$
- ▶ To visualise the k -means clusters, I am using the two principal components (PCA) of \mathbf{X} – not great in this example because they explain $\sim 70\%$ of the data's variance
- ▶ So, actually clustering applied on a matrix $\mathbf{Z} \in \mathbb{R}^{2 \times 150}$
- ▶ Download data: dropbox.com/s/rgyg190whw26qrj/data-COMP0084-intro-to-ml.zip?dl=0


Clustering — k -means, an example

- ▶ $k = 2$
- ▶ clusters are denoted by C_i
- ▶ a cross is used to denote each cluster's centroid




Clustering — k -means, an example

- ▶ $k = 2$
- ▶ clusters are denoted by C_i
- ▶ a cross is used to denote each cluster's centroid
- ▶ which search queries are closer to their cluster's centroid?




Clustering — k -means, an example

- ▶ $k = 3$
- ▶ clusters are denoted by C_i
- ▶ a cross is used to denote each cluster's centroid
- ▶ which search queries are closer to their cluster's centroid?




Clustering – k -means, an example

- ▶ $k = 3$
- ▶ clusters are denoted by C_i
- ▶ a cross is used to denote each cluster's centroid
- ▶ which search queries are closer to their cluster's centroid?
- ▶ does the addition of a cluster change the thematic coverage of the revised clusters?
- ▶ central queries have changed




Clustering – k -means, an example

- ▶ $k = 4$
- ▶ clusters are denoted by C_i
- ▶ a cross is used to denote each cluster's centroid
- ▶ which search queries are closer to their cluster's centroid?
- ▶ does the addition of a cluster change the thematic coverage of the revised clusters?
- ▶ central queries have changed



Clustering — k -means, an example

- ▶ $k = 5$
- ▶ clusters are denoted by C_i
- ▶ a cross is used to denote each cluster's centroid
- ▶ which search queries are closer to their cluster's centroid?
- ▶ does the addition of a cluster change the thematic coverage of the revised clusters?
- ▶ central queries have changed — *partially!*



Next lectures with me

Topic models and vector semantics (*word embeddings*)

- ▶ March 2 and 3 (2 hours + 1 hour guest lecture)

Guest lecture (*self-invited*) about modelling infectious disease prevalence (*including COVID-19*) using web search activity

- ▶ March 16 (1 hour)